Lecture 2:
Data analysis

Universitat de les Illes Balears
Departament de Ciències Matemàtiques i Informàtica

11752 Aprendizaje Automático
11752 Machine Learning
Máster Universitario
en Sistemas Inteligentes

Alberto ORTIZ RODRÍGUEZ

Contents

- **Feature engineering**
  - *Given a machine learning problem and a collection of features, select the **minimum set** (and/or maybe a **transformation**) that improves/retains as much as possible the ability to discriminate among samples*
    - Generally we find a large number of possible features
      - the more features, the more complex the classifier is
        more parameters, harder to tune $\Rightarrow$ **feature selection** / **dimensionality reduction**
      - the computational complexity related to their calculation can be taken into account
    - Consider whether a **transformation** of the features could enhance performance
    - Given the number of samples N, the number of features L should be:
      - big enough to learn
        what makes classes different
        what makes individuals of the same class similar
      - small enough not to make individuals of the same class different
      - it is well known that the **classification error** gets lower as the ratio N/L gets higher
      - as a **rule of thumb**, $N/L > 3$, though, in some cases, $N/L > 10 \text{ or } 20$
      - in other words, the more features, the more samples you need

- **Feature engineering**
  - **General criterion:**
    - select those features that result in a large between-class distance and a reduced variance between class elements (within-class variance)

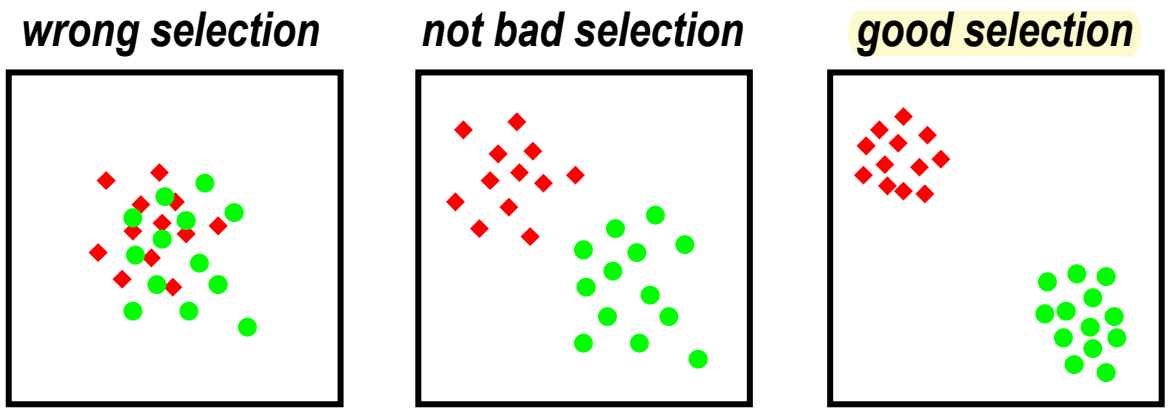

  - **Actions to do** to "engineer" the feature set:
    - understand your data, i.e. explore your data (maybe to know which is your case above)
    - if needed, transform the data to compensate well known problems during training
    - examine features in isolation
    - examine features in combination
    - combine your features

    → feature selection / dimensionality reduction

Contents

Data exploration

- **Data exploration** is the first, basic step to understand your data
- This includes **data visualization** for **qualitative assessment,** detection of **anomalies**, **trends** and **relationships**, as well as to detect the necessity for **data cleaning**
- Let us consider the **Iris** flower dataset (Fisher's Iris data set)
  - multivariate dataset by the British statistician and biologist Ronald Fisher (1936)
  - 150 samples under four attributes:
    - sepal length
    - sepal width
    - petal length
    - petal width
  - 3 species:
    - setosa
    - versicolor
    - virginica

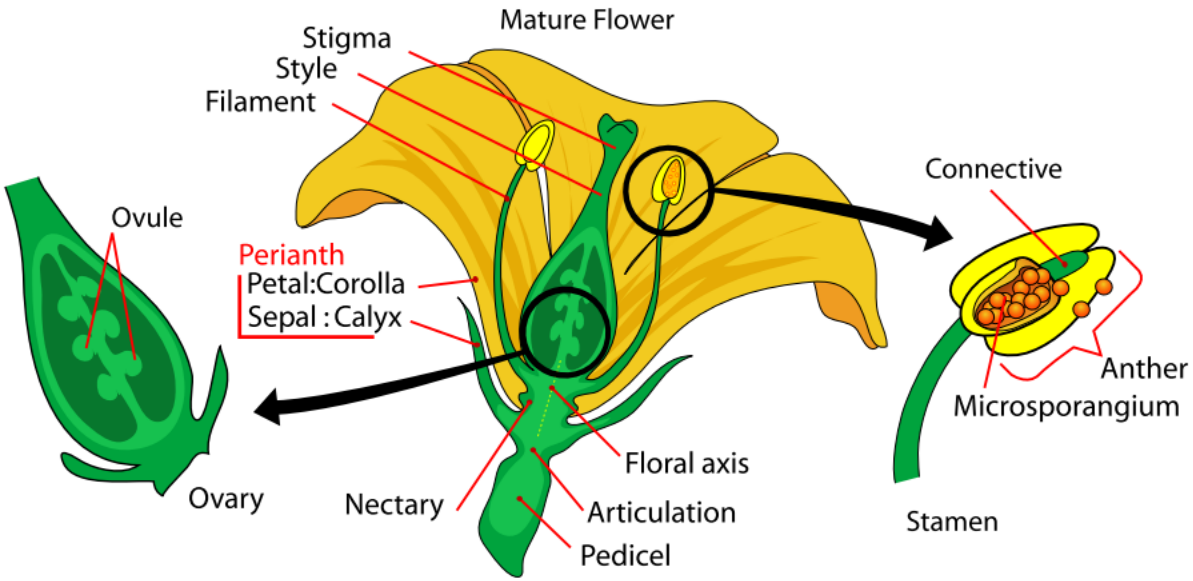

- Basic **descriptive data**:

```
import numpy as np
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
y = iris.target
nos, nod = X.shape
print('no. samples = %d, no. dimensions = %d' % (nos, nod))
noc = len(np.unique(y))
print('no. classes = %d' % (noc))
mu = np.mean(X, axis=0)
std = np.std(X, axis=0)
std2 = np.var(X, axis=0)
print('    mean std  var')
for i in range(nod):
    print('x%d: %.2f %.2f %.2f' % (i+1,mu[i],std[i],std2[i]))
```

```
no. samples = 150, no. dimensions = 4
no. classes = 3
    mean std  var
x1: 5.84 0.83 0.68
x2: 3.06 0.43 0.19
x3: 3.76 1.76 3.10
x4: 1.20 0.76 0.58
```

Data exploration

- Basic **descriptive data**:

```
import numpy as np
from sklearn.datasets import load_iris
iris = load_iris()
print(iris.DESCR)
```

```
Iris plants dataset
--------------------

**Data Set Characteristics:**

    :Number of Instances: 150 (50 in each of three classes)
    :Number of Attributes: 4 numeric, predictive attributes and the class
    :Attribute Information:
        - sepal length in cm
        - sepal width in cm
        - petal length in cm
        - petal width in cm
        - class:
                - Iris-Setosa
                - Iris-Versicolour
                - Iris-Virginica

    :Summary Statistics:

    ============== ==== ==== ======= ===== ====================
                    Min  Max   Mean    SD   Class Correlation
    ============== ==== ==== ======= ===== ====================
    sepal length:   4.3  7.9   5.84   0.83    0.7826
    sepal width:    2.0  4.4   3.05   0.43   -0.4194
    petal length:   1.0  6.9   3.76   1.76    0.9490  (high!)
    petal width:    0.1  2.5   1.20   0.76    0.9565  (high!)
    ============== ==== ==== ======= ===== ====================

    :Missing Attribute Values: None
    :Class Distribution: 33.3% for each of 3 classes.
```

Data exploration

- Basic **visualization**:

```
import matplotlib.pyplot as plt
plt.figure()
for c in range(noc):
    i = np.where(y == c)[0]
    plt.scatter(X[i,0],X[i,1])
plt.xlabel('x1 (sepal length)')
plt.ylabel('x2 (sepal width)')
plt.title('iris dataset')
plt.show()
```

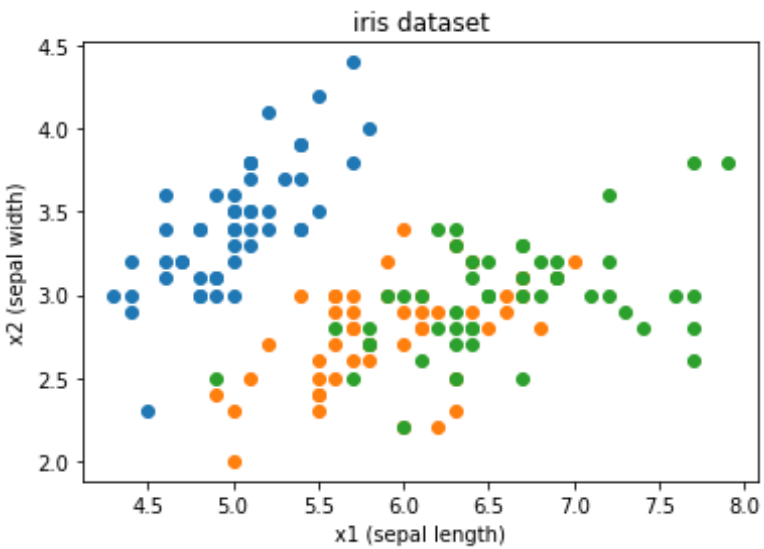

- For multidimensional datasets, i.e. more than 2 / 3 dimensions, the standard methods of visualization are not an option
  - Among many others:
    - the **Scatter Plot Matrix (SPLOM)** and
    - the **parallel coordinates plot**

  are alternative visualization tools, though of limited capability

Data exploration

- **Scatter PLOt Matrix**
  - Correlation plots & Histograms

```
import seaborn as sb
df = sb.load_dataset('iris')
sb.pairplot(df, hue='species')
```

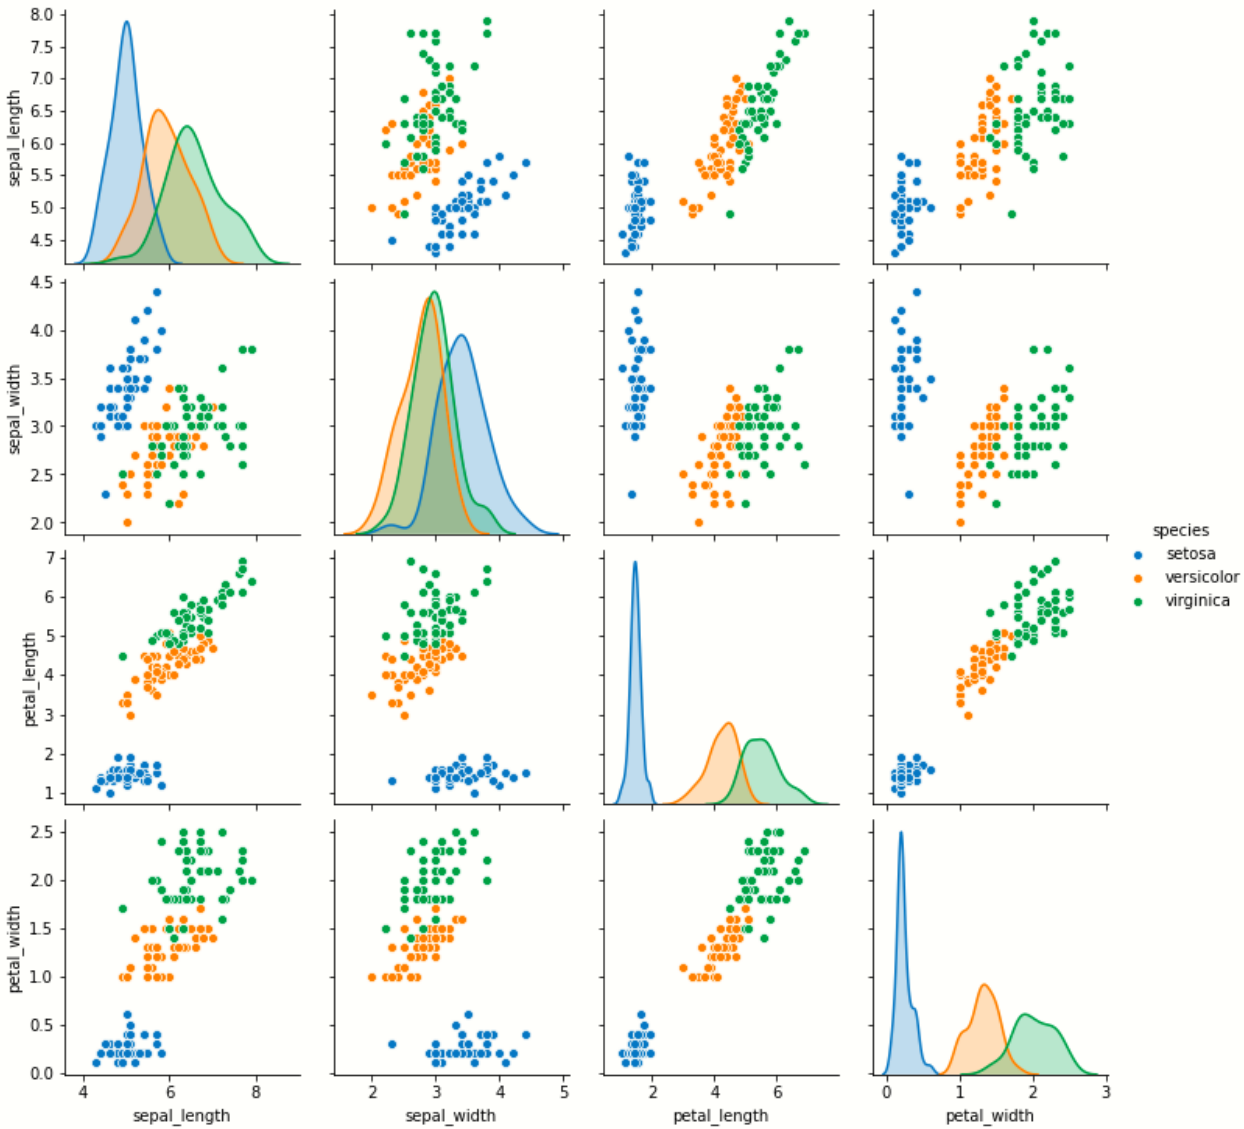

number of features

- **Parallel coordinates plot**

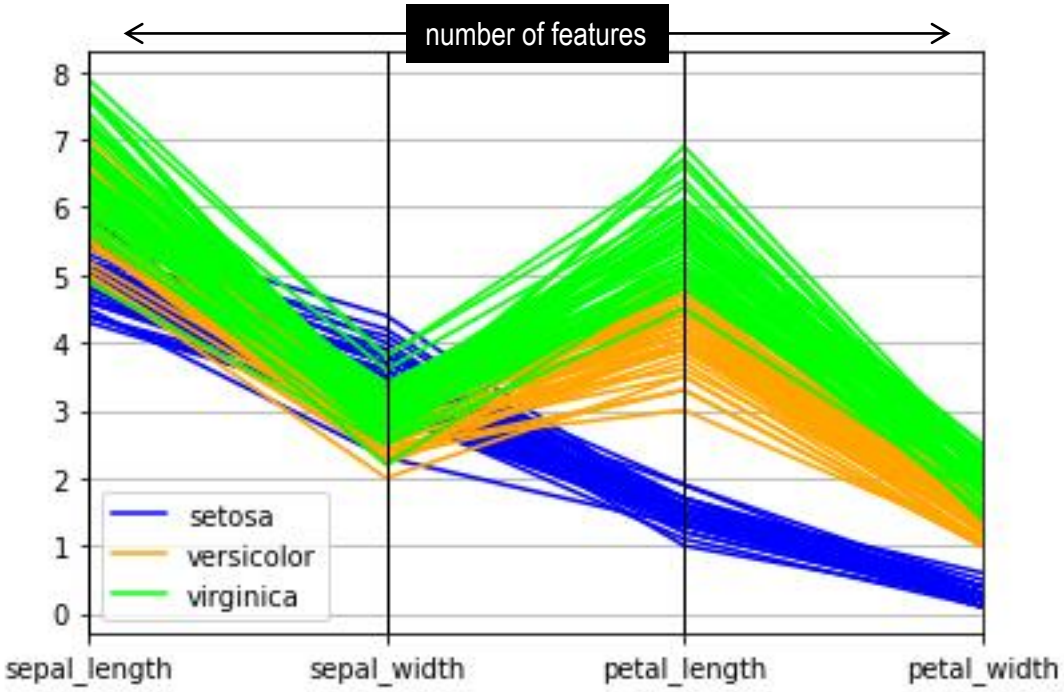


```
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sb
df = sb.load_dataset('iris')
pd.plotting.parallel_coordinates(df, 'species', color=('#0000FF', '#FFA500', '#00FF00'))
plt.legend(loc='lower left')
plt.show()
```

Data exploration

- **Pandas** is a library for data manipulation and analysis which can be useful for ML
  - The pandas **dataframe class** may be particularly useful for data manipulation and indexing, as well as for file input / output
- To create a *dataframe* we need an array of values.
  Moreover, we can add labels for the columns and for the samples:

```
import pandas as pd
df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
     index=['cobra', 'viper', 'sidewinder'],
     columns=['max_speed', 'shield'])
print(df.head())
```

```
             max_speed  shield
cobra                1       2
viper                4       5
sidewinder           7       8
```

- In *dataframes*, indexing can be very flexible with the **df.loc()** method:

```
print(df.loc[['viper', 'sidewinder']])
print(df.loc['cobra':'viper', 'max_speed'])
print(df.loc[df['shield'] > 4, ['max_speed']])
```

```
             max_speed  shield
viper                4       5
sidewinder           7       8

cobra    1
viper    4
Name: max_speed, dtype: int64

             max_speed
viper                4
sidewinder           7
```

Data exploration

- Let us use the *Titanic* dataset to illustrate other functionalities of *dataframes*:

```
import seaborn as sb
titanic = sb.load_dataset('titanic')
df = titanic
print(df.info())
print(df.head(3))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column       Non-Null Count  Dtype
---  ------       --------------  -----
 0   survived     891 non-null    int64
 1   pclass       891 non-null    int64
 2   sex          891 non-null    object
 3   age          714 non-null    float64
 4   sibsp        891 non-null    int64
 5   parch        891 non-null    int64
 6   fare         891 non-null    float64
 7   embarked     889 non-null    object
 8   class        891 non-null    category
 9   who          891 non-null    object
 10  adult_male   891 non-null    bool
 11  deck         203 non-null    category
 12  embark_town  889 non-null    object
 13  alive        891 non-null    object
 14  alone        891 non-null    bool
dtypes: bool(2), category(2), float64(2),
int64(4), object(5)
memory usage: 80.7+ KB
```

```
 survived pclass    sex  age ... deck embark_town alive alone
0        0      3   male 22.0 ...  NaN Southampton    no Fale
1        1      1 female 38.0 ...    C   Cherbourg   yes Fale
2        1      3 female 26.0 ...  NaN Southampton   yes  True
```

  - **df.tail(n)** displays the last *n* samples

- We can also load the dataset from disk. Let us assume the dataset is in file *titanic.csv*:

```
df = pd.read_csv('titanic.csv')
print(df.info())
```

  - Other formats also available for input/output, e.g. JSON, excel, etc.

Data exploration

- The **df.describe()** method provides a summary of the dataset statistics:

```
print(df.describe())
```

```
          survived      pclass         age       sibsp       parch        fare
count  891.000000  891.000000  714.000000  891.000000  891.000000  891.000000
mean     0.383838    2.308642   29.699118    0.523008    0.381594   32.204208
std      0.486592    0.836071   14.526497    1.102743    0.806057   49.693429
min      0.000000    1.000000    0.420000    0.000000    0.000000    0.000000
25%      0.000000    2.000000   20.125000    0.000000    0.000000    7.910400
50%      0.000000    3.000000   28.000000    0.000000    0.000000   14.454200
75%      1.000000    3.000000   38.000000    1.000000    0.000000   31.000000
max      1.000000    3.000000   80.000000    8.000000    6.000000  512.329200
```

- For selecting elements of the dataset, one can additionally use column labels and the **df.iloc()** method:

```
X1 = df.iloc[:,[1,2,3,4,5,6]].to_numpy()
X2 =
df[['pclass','sex','age','sibsp','parch','fare']]
y  = df['survived']
print(X1[0:3,:])
print(X2.head(3))
```

```
[[3 'male'   22.0 1 0 7.25]
 [1 'female' 38.0 1 0 71.2833]
 [3 'female' 26.0 0 0 7.925]]

  pclass     sex  age sibsp parch     fare
0      3    male 22.0     1     0   7.2500
1      1  female 38.0     1     0  71.2833
2      3  female 26.0     0     0   7.9250
```

Data exploration

- Conditions can also be used for selecting samples:

```
print(df[df['deck'] == 'C'].head(3))
```

```
urvived pclass    sex  age ... deck embark_town alive alone
1        1      1 female 38.0 ...    C   Cherbourg   yes False
3        1      1 female 35.0 ...    C Southampton   yes False
11       1      1 female 58.0 ...    C Southampton   yes  True
```

```
print(df[(df['age'] > 50) & (df['pclass'] < 2)].head(3))
```

```
 survived pclass    sex  age ... deck embark_town alive alone
6        0      1   male 54.0 ...    E Southampton    no  True
11       1      1 female 58.0 ...    C Southampton   yes  True
54       0      1   male 65.0 ...    B   Cherbourg    no False
```

- The *dataframe* object provides a number of ways to get more details of the dataset:
  - The **df.columns** attribute is a list with the labels of the dataset columns
  - **df.values()** or **df.to_numpy()** provide the dataset values as a numpy array
  - **df.count_values()** returns the number of times the different values occur in a column
  - With **df.nunique()** we can see the counts of unique values in each column

```
print(df[['age','deck']].nunique())
print(df['sex'].value_counts())
```

```
age       88
deck       7

male       577
female     314
Name: sex, dtype: int64
```

Data exploration

- We can **remove some features** (columns) which are useless:

```
udf = df
udf.drop('embarked',axis=1,inplace=True)
udf.drop('class',axis=1,inplace=True)
udf.drop('who',axis=1,inplace=True)
udf.drop('adult_male',axis=1,inplace=True)
udf.drop('deck',axis=1,inplace=True)
udf.drop('embark_town',axis=1,inplace=True)
udf.drop('alive',axis=1,inplace=True)
udf.drop('alone',axis=1,inplace=True)
print(udf.info())
```

```
RangeIndex: 891 entries, 0 to 890
Data columns (total 7 columns):
 #   Column     Non-Null Count  Dtype
---  ------     --------------  -----
 0   survived   891 non-null    int64
 1   pclass     891 non-null    int64
 2   sex        891 non-null    object
 3   age        714 non-null    float64
 4   sibsp      891 non-null    int64
 5   parch      891 non-null    int64
 6   fare       891 non-null    float64
```

- We can as well **drop duplicates**, if any:

```
import pandas as pd
df = pd.DataFrame({
    'brand': ['Yum','Yum','Indo','Indo','Indo'],
    'style': ['cup','cup','cup','pack','pack'],
    'rating': [4, 4, 3.5, 15, 5]
})
print(df)
print(df.drop_duplicates())
print(df.drop_duplicates(subset='brand'))
```

```
   brand style  rating
0    Yum   cup     4.0
1    Yum   cup     4.0
2   Indo   cup     3.5
3   Indo  pack    15.0
4   Indo  pack     5.0

   brand style  rating
0    Yum   cup     4.0
2   Indo   cup     3.5
3   Indo  pack    15.0
4   Indo  pack     5.0

   brand style  rating
0    Yum   cup     4.0
2   Indo   cup     3.5
```

Contents

Data preprocessing

- Preparation of data samples before proceeding to their use
  - **Handling categorical data**
  - **Outlier detection (and removal)**
  - **Data normalization / standardization**
  - **Filling in missing data**

Data preprocessing: Categorical data

- **Handling categorical data**
  - Categorical data must be converted to **numeric values** before learning
    - The **LabelEncoder()** object assigns a **progressive integer label** to every class label

```
import seaborn as sb
from sklearn.preprocessing import LabelEncoder
titanic = sb.load_dataset('titanic')
df = titanic
print(df['sex'][:5])
le = LabelEncoder()
df['sex'] = le.fit_transform(df['sex'])
print(df['sex'][:5])
print(le.classes_)
```

```
0      male
1    female
2    female
3    female
4      male
Name: Sex, dtype: object
0    1
1    0
2    0
3    0
4    1
Name: Sex, dtype: int32
['female' 'male']
```

    - The names of the classes are in attribute *le.classes_*

  - Unfortunately, on some occasions, this is not a good encoding for training, and ***one-hot encoding*** must be used instead:

```
from sklearn.preprocessing import OneHotEncoder
df = titanic
ohe = OneHotEncoder()
data = np.expand_dims(df['sex'], axis=-1)
ohe.fit(data)
data_ = ohe.transform(data).toarray()
print(data_[:5])
```

```
[[0. 1.]
 [1. 0.]
 [1. 0.]
 [1. 0.]
 [0. 1.]]
```

Data preprocessing: Outlier detection

- **Outlier detection (and removal)**
  - ***outlier*** ≡ sample that does not agree with the rest of the population

- normally, distance to the mean is $k\sigma$, $k\uparrow\uparrow$
- an outlier can distort training
  - the resulting classifier / regressor may not classify / predict for new samples in the right way

e.g. sonar readings (case 1D)

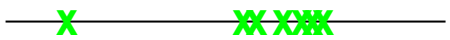

e.g. case 2D

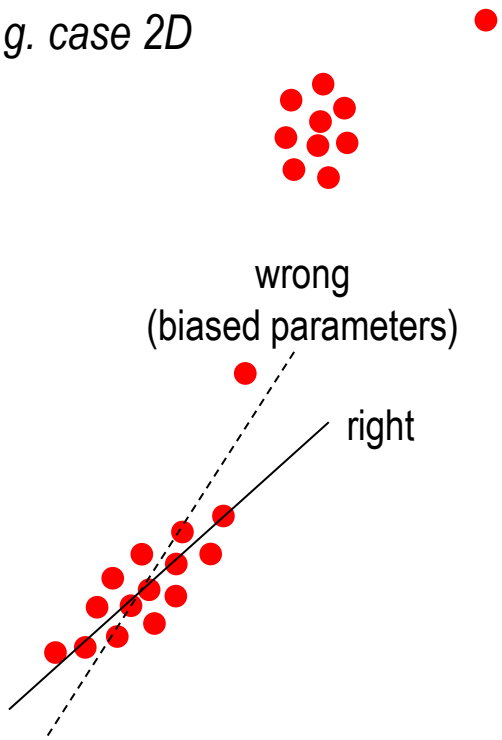

Data preprocessing: Outlier detection

- **Sources of outliers**:
  - measurement error (instrument error) or experimental error (wrong data extraction)
  - data entry error (data collection/typing) or data processing error
  - extreme noisy sample (natural outlier)
- If you need to counteract the outliers, these are some of the **possible actions**:
  - **discard** the outliers (= full sample) if the dataset permits to do so, i.e. it is big enough
  - alter the data:
    - **trimming**: extreme values are set to "missing", i.e. NaN (Python)
    - **winsorization**: replace given parts of a sample at the high and low ends with the most extreme values, i.e. use $k\sigma$, e.g. $k = 3$

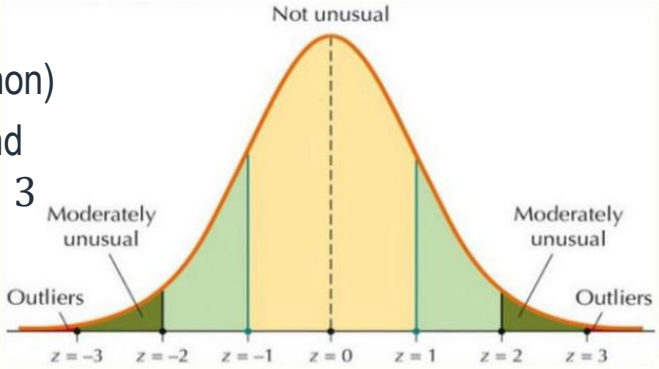

Winsorized mean. After sorting the data, we replace x_1 and x_{10} by resp. x_2 and x_9

$$\frac{\overbrace{x_2 + x_2} + x_3 + x_4 + x_5 + x_6 + x_7 + x_8 + \overbrace{x_9 + x_9}}{10}$$

(20% winsorized mean)

1, 5, 7, 8, 9, 10, 10, 12, 12, 34 → μ = 10.8
5, 5, 7, 8, 9, 10, 10, 12, 12, 12 → μ = 9.0

scipy.stats.mstats.winsorize()

  - **tolerate** the outliers by reducing their influence
    - use optimization methods from **robust statistics** (large values are attenuated "on-line")

Data preprocessing: Outlier detection

- **z-score method (Gausssian data)**: *$\mu \pm 3\sigma$ accumulates 99.7% of the probability*

```
import numpy as np
from sklearn.datasets import load_wine
wine = load_wine()
X = wine.data
y = wine.target
print(wine.DESCR)

outl = np.zeros((3,13))
for c in range(3):
    for f in range(13):
        cc = X[y == c, f]
        mu, sg = np.mean(cc), np.std(cc)
        cut_off = sg * 3
        lower, upper = mu - cut_off, mu + cut_off
        # identify outliers
        outliers = [x for x in cc if x < lower or x > upper]
        nout = len(outliers)
        # non-outliers
        non_outliers = [x for x in cc if x >= lower and x <= upper]
        nok = len(non_outliers)
        outl[c,f] = nout
print(outl)
```

```
from scipy import stats
z = np.abs(stats.zscore(X))
# discard samples with z > 3
```

$$z = \frac{x - \mu}{\sigma}$$

```
============================ ==== ===== ======= =====
                              Min   Max   Mean     SD
============================ ==== ===== ======= =====
Alcohol:                     11.0  14.8    13.0    0.8
Malic Acid:                  0.74  5.80    2.34   1.12
Ash:                         1.36  3.23    2.36   0.27
Alcalinity of Ash:           10.6  30.0    19.5    3.3
Magnesium:                   70.0 162.0    99.7   14.3
Total Phenols:               0.98  3.88    2.29   0.63
Flavanoids:                  0.34  5.08    2.03   1.00
Nonflavanoid Phenols:        0.13  0.66    0.36   0.12
Proanthocyanins:             0.41  3.58    1.59   0.57
Colour Intensity:             1.3  13.0     5.1    2.3
Hue:                         0.48  1.71    0.96   0.23
OD280/OD315 of diluted wines: 1.27  4.00    2.61   0.71
Proline:                      278  1680     746    315
============================ ==== ===== ======= =====

:Missing Attribute Values: None
:Class Distribution: class_0 (59), class_1 (71), class_2 (48)
```

class	f1	f2	f3	f4	f5	f6	f7	f8	f9	f10	f11	f12	f13
1	0	0	1	1	0	2	0	1	0	0	0	0	0
2	0	1	1	0	2	0	1	0	1	1	1	0	0
3	0	0	0	0	0	1	0	0	1	0	0	0	0

With this code, we know that there are outliers in all classes, but we should discover which samples are affected !!

and next do trimming or winsorizing

Data preprocessing: Outlier detection

- **Inter-quartile range method (non-Gausssian data)**

```
import numpy as np
from sklearn.datasets import load_wine
wine = load_wine()
X = wine.data
y = wine.target

from numpy import percentile
outl = np.zeros((3,13))
for c in range(3):
    for f in range(13):
        cc = X[y == c, f]
        q25, q75 = percentile(cc, 25), percentile(cc, 75)
        iqr = q75 - q25
        cut_off = iqr * 1.5
        lower, upper = q25 - cut_off, q75 + cut_off
        # identify outliers
        outliers = [x for x in cc if x < lower or x > upper]
        nout = len(outliers)
        # non-outliers
        non_outliers = [x for x in cc if x >= lower and x <= upper]
        nok = len(non_outliers)
        outl[c,f] = nout
print(outl)
```

- IQR = difference between the 75th and the 25th percentiles of the data (Q3, Q1)
- Situates outliers out of the $\pm k \times \mathrm{IQR}$ interval

class	f1	f2	f3	f4	f5	f6	f7	f8	f9	f10	f11	f12	f13
1	0	9	1	3	0	2	0	4	4	1	0	0	0
2	3	7	2	4	5	0	1	0	8	4	1	0	1
3	0	0	0	0	0	2	1	1	2	0	0	2	0

With this code, we know that there are outliers in all classes, but we should discover which samples are affected !!

and next do trimming or winsorizing

Data preprocessing: Outlier detection

- **Box-plots**

```
import matplotlib.pyplot as plt
import seaborn as sb        # also in matplotlib
df = sb.load_dataset('iris')
plt.figure()
sb.boxplot(y=df['species'], x=df['sepal_length'])
plt.show()
plt.figure()
sb.boxplot(y=df['species'], x=df['sepal_width'])
plt.show()
plt.figure()
sb.boxplot(y=df['species'], x=df['petal_length'])
plt.show()
plt.figure()
sb.boxplot(y=df['species'], x=df['petal_width'])
plt.show()
```

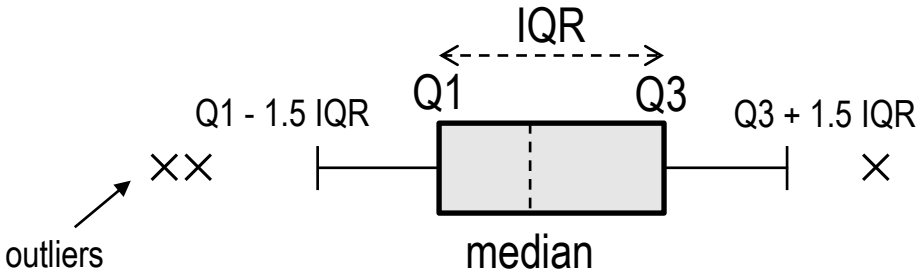

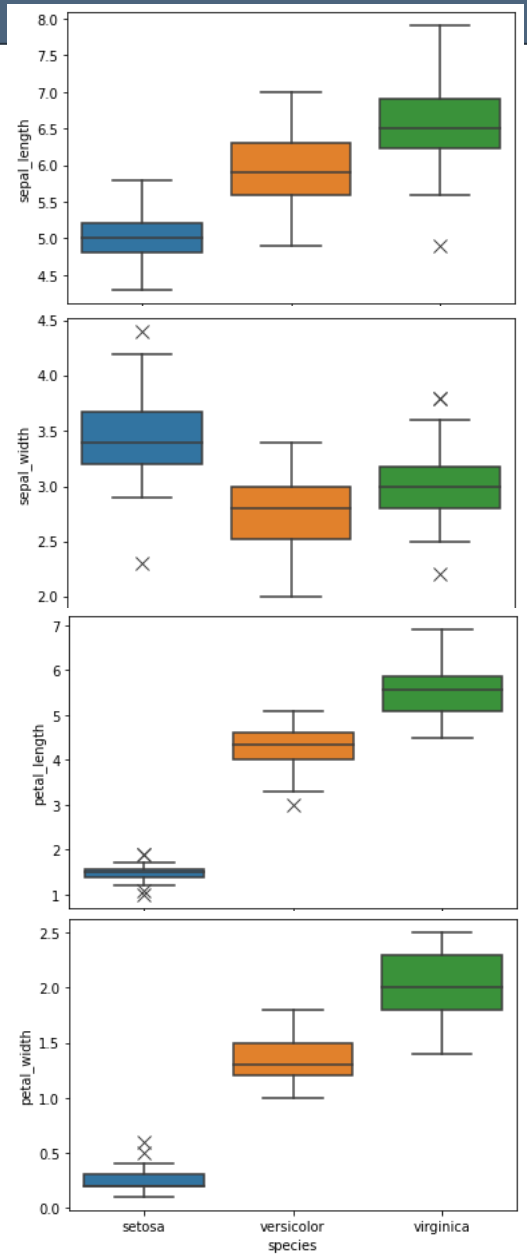

Data preprocessing: Outlier detection

- **Box-plots**
  - Why 1.5IQR?
    - Related with the 68–95–99 rule from the Gaussian distribution
    - In the Gaussian distribution, ±1.5IQR covers approx. the same probability as ±3σ

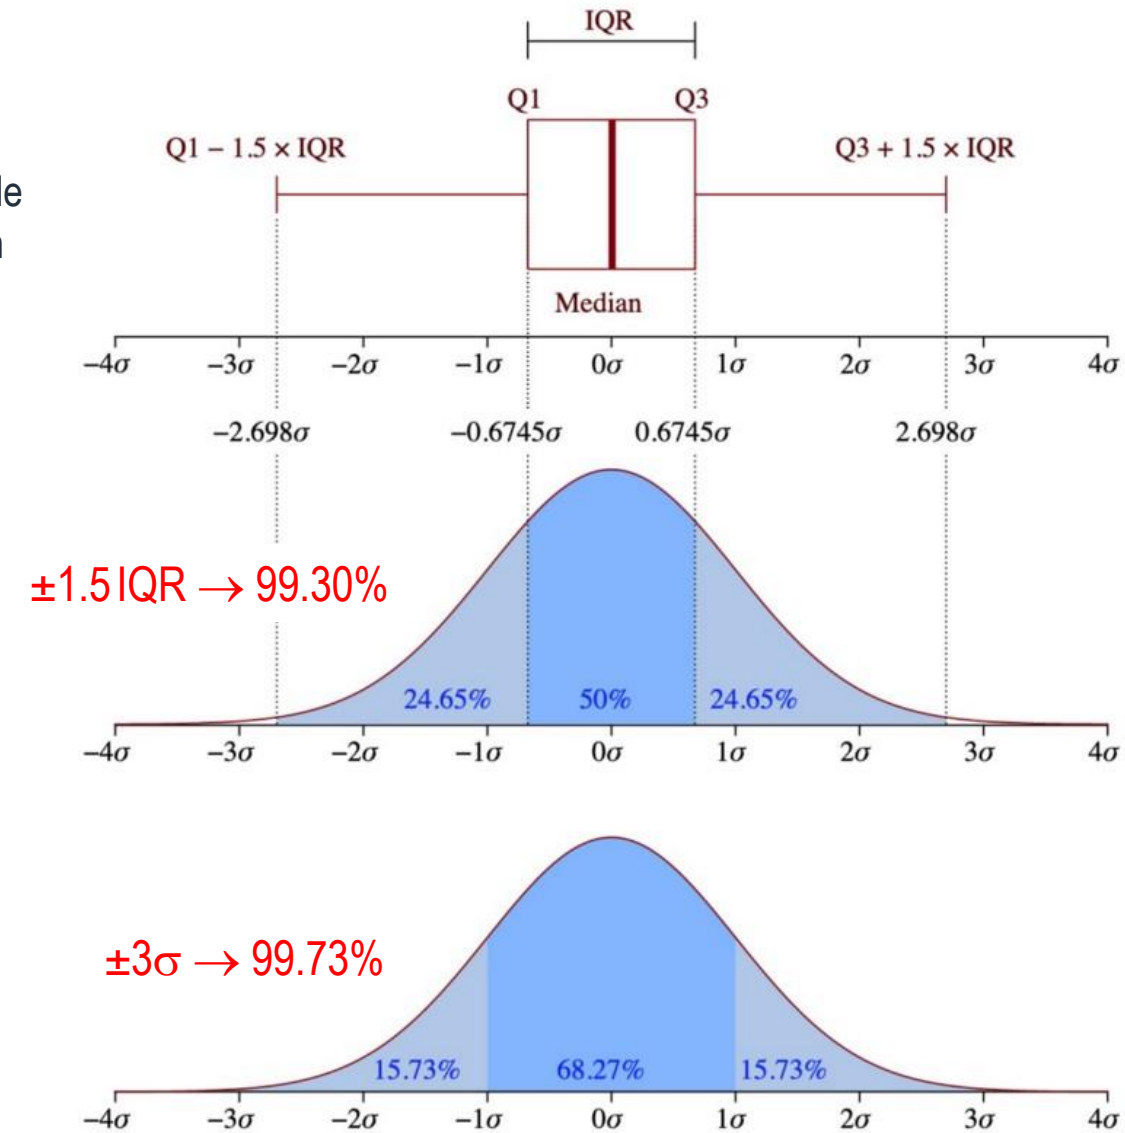

Data preprocessing: Outlier detection

- **Automatic detection of outliers**
  - **Local Outlier Factor (LOF):** measures the local deviation of the density of a sample with respect to its neighbors

```
import numpy as np
import pandas as pd
from sklearn.neighbors import LocalOutlierFactor
import matplotlib.pyplot as plt
import seaborn as sb
df = sb.load_dataset('iris')
data = df.values
X = data[:,:-1]
y = data[:,-1]
# identify outliers in the training dataset
lof = LocalOutlierFactor(n_neighbors=20)
yhat = lof.fit_predict(X)
# select all rows that are not outliers
mask = yhat != -1
X, y = X[mask, :], y[mask]
y = np.expand_dims(y,axis=1)
df2 = pd.DataFrame(np.hstack((X,y)))
df2.columns = df.columns
```

  - Others: **IsolationForest**, etc.

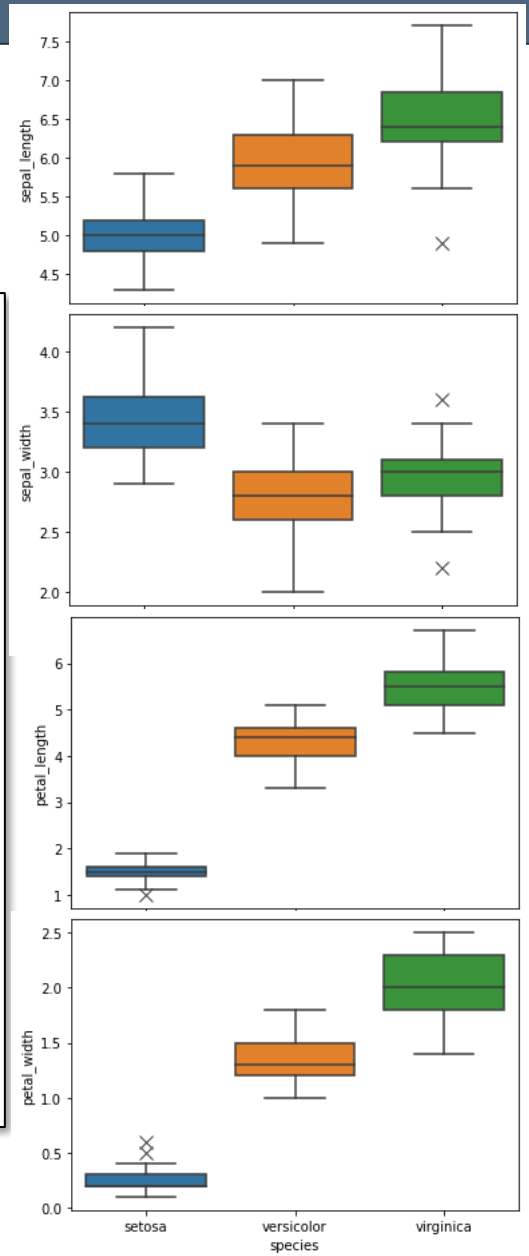

- **Data normalization / standardization**
  - Often, different features do not have the same **dynamic range** (range of values)
    - characteristics with wider ranges will have **more influence on the classification** regardless of whether they are more relevant to the design of the classifier or not

  - Solution:
    - normalize/scale features so that their dynamic ranges are similar
      - **linear scaling**
        - mu-sigma normalization (standardization)
        - max-min normalization
        - others
      - **non-linear scaling**
        - *softmax* normalization
        - others

Data preprocessing: Normalization

- **μ-σ (mu-sigma) normalization**
  - Given L – feature descriptors:

$$\forall k = 1, \ldots, L, \quad \overline{x}_k = \frac{\sum_i x_{ik}}{N}$$

$$\sigma_k^2 = \frac{\sum_i (x_{ik} - \overline{x}_k)^2}{N-1}$$

$$\widehat{x}_{ik} = \frac{x_{ik} - \overline{x}_k}{\sigma_k}$$

  - After the transformation:

$$\mathrm{E}[\widehat{x}_{ik}] = 0, \quad \mathrm{Var}[\widehat{x}_{ik}] = 1$$

$$\boxed{\begin{array}{lcl} x_{ik} - \overline{x}_k = 0 & \Rightarrow & \widehat{x}_{ik} = 0 \\ x_{ik} - \overline{x}_k = +k\sigma & \Rightarrow & \widehat{x}_{ik} = +k \\ x_{ik} - \overline{x}_k = -k\sigma & \Rightarrow & \widehat{x}_{ik} = -k \end{array}}$$

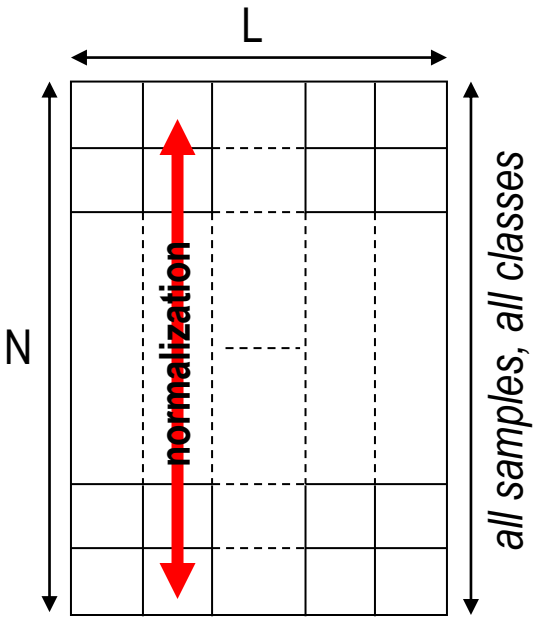

VERY IMPORTANT

Apply the transformation to the full dataset once.

- **Typical mistake**. Standardize the training set separately from the test set.
- Keep the transformation parameters $(\bar{x}, \sigma)$ to standardize new samples

Data preprocessing: Normalization

- **μ-σ (mu-sigma) normalization**

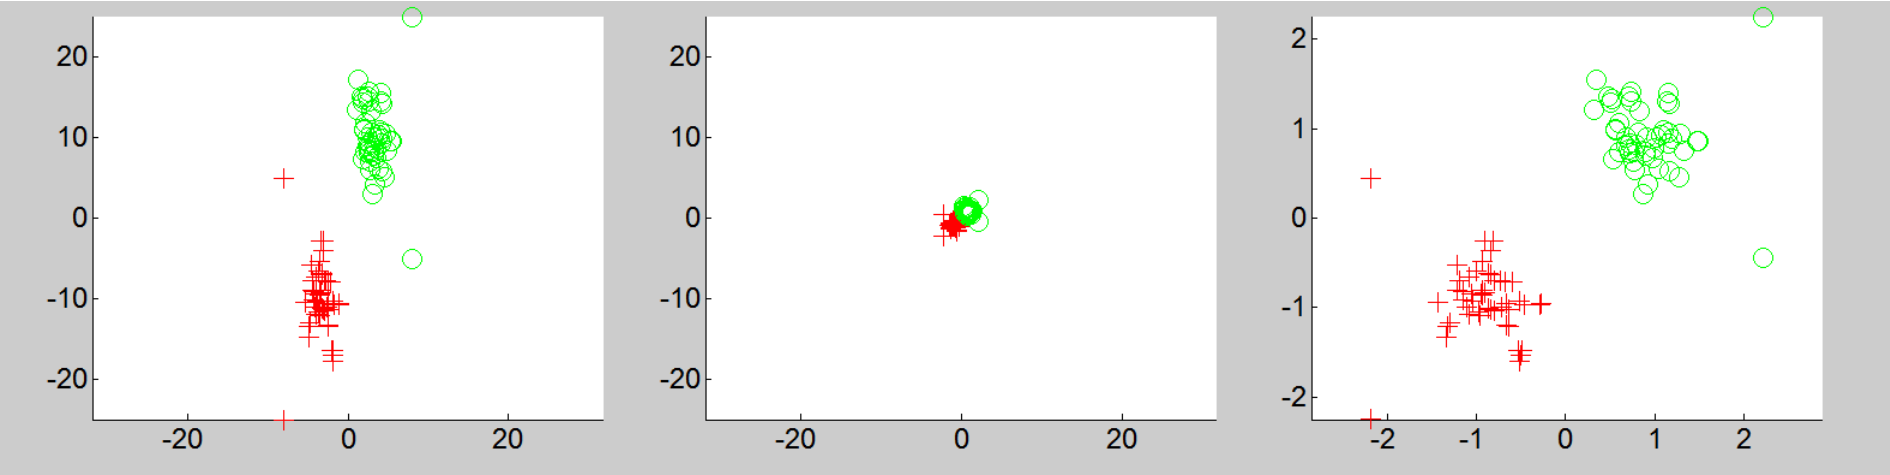


```
normalization: mu-sigma
k=1 org:  -8.00 -   8.00 :  16.00
k=2 org: -25.00 -  25.00 :  50.00
ratio  :   3.13
k=1 nor:  -2.18 -   2.21 :   4.38
k=2 nor:  -2.25 -   2.25 :   4.50
ratio  :   1.03
```

← dynamic range of x_1

← dynamic range of x_2

← dynamic range of $\hat{x}_1$

← dynamic range of $\hat{x}_2$

Data preprocessing: Normalization

- **Max-min normalization**
  - Given L – feature descriptors:

$$\forall k = 1, \ldots, L, \quad X_k = \max_i \{x_{ik}\}$$
$$x_k = \min_i \{x_{ik}\}$$
$$\widehat{x}_{ik} = \frac{x_{ik} - x_k}{X_k - x_k}$$

  - After the transformation: $\widehat{x}_{ik} \in [0, 1]$

$$\boxed{\begin{aligned} x_{ik} = \max_i \{x_{ik}\} &\Rightarrow \widehat{x}_{ik} = 1 \\ x_{ik} = \min_i \{x_{ik}\} &\Rightarrow \widehat{x}_{ik} = 0 \end{aligned}}$$

... distributes the data within the range [0,1]

- the original ends correspond to 0 and 1
- e.g. if originally the range of values was [-30, 100], after normalization, value -30 will become 0 for that feature, while a value of 100 will become 1

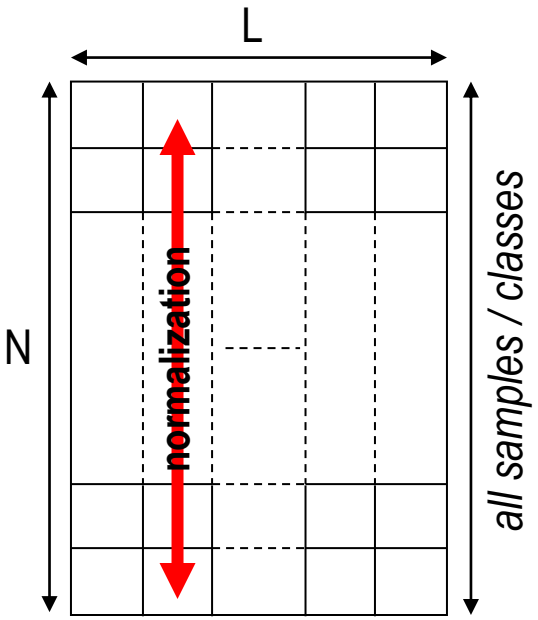

VERY IMPORTANT:
Apply the transformation to the full dataset once.

❖ Keep the transformation parameters (x, X) to normalize new samples

Data preprocessing: Normalization

- **Max-min normalization**

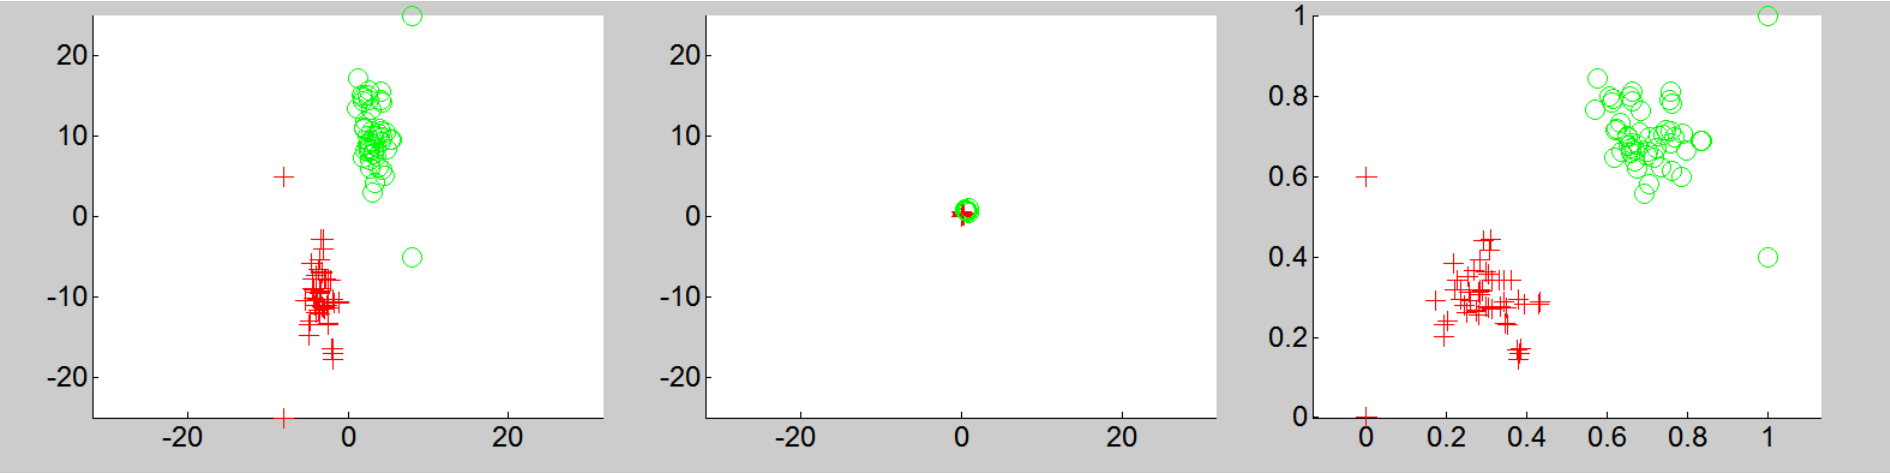


```
normalization: min-max
k=1 org:  -8.00 -   8.00 :  16.00
k=2 org: -25.00 -  25.00 :  50.00
ratio  :   3.13
k=1 nor:   0.00 -   1.00 :   1.00
k=2 nor:   0.00 -   1.00 :   1.00
ratio  :   1.00
```

← dynamic range of x_1

← dynamic range of x_2

← dynamic range of $\hat{x}_1$

← dynamic range of $\hat{x}_2$

Data preprocessing: Normalization

- ***Softmax* normalization** (non-linear transformation)
  - Given L – feature descriptors:

$$\forall k = 1, \ldots, L, \quad \overline{x}_k = \frac{\sum_i x_{ik}}{N}$$

$$\sigma_k^2 = \frac{\sum_i (x_{ik} - \overline{x}_k)^2}{N-1}$$

$$z_{ik} = \frac{x_{ik} - \overline{x}_k}{r\sigma_k}$$

$$\widehat{x}_{ik} = \frac{1}{1 + e^{-z_{ik}}}$$

  - After the transformation:

$$\widehat{x}_{ik} \in [0, 1]$$

$$\boxed{\begin{aligned} x_{ik} = \overline{x}_k &\Rightarrow \widehat{x}_{ik} = \frac{1}{2} \\ x_{ik} = +\infty &\Rightarrow \widehat{x}_{ik} = 1 \\ x_{ik} = -\infty &\Rightarrow \widehat{x}_{ik} = 0 \end{aligned}}$$

… but it does not distribute evenly the data within [0,1]

- exponentially "concentrates" values far from the mean as a function of σ and r:
  - the higher r, the closer get the farthest samples

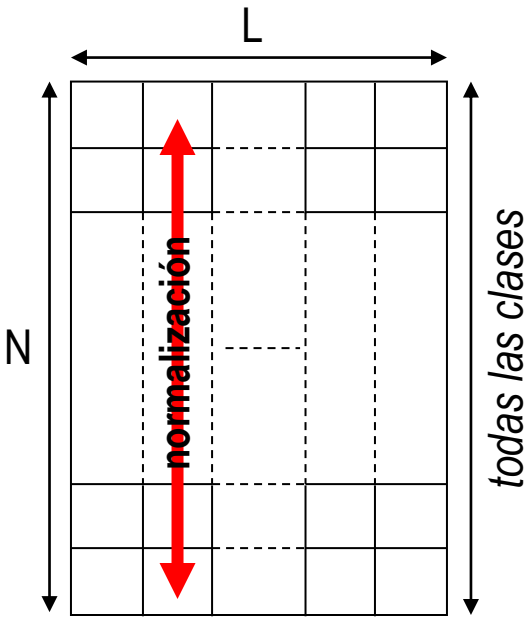

VERY IMPORTANT:
Apply the transformation to the full dataset once.

❖ Keep the transformation parameters $(\bar{x}, \sigma, r)$ to standardize new samples

- ***Softmax* normalization**

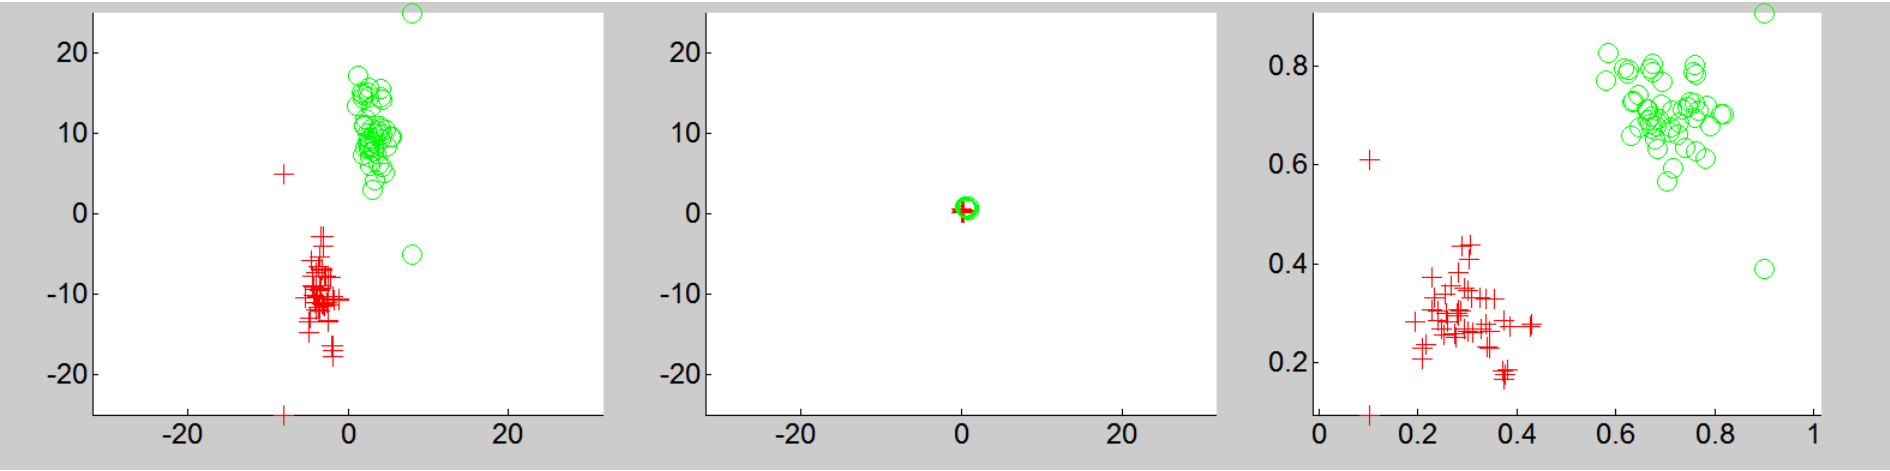


```
normalization: softmax (r=1)
k=1 org:  -8.00 -   8.00 :  16.00
k=2 org: -25.00 -  25.00 :  50.00
ratio  :   3.13
k=1 nor:   0.10 -   0.90 :   0.80
k=2 nor:   0.10 -   0.90 :   0.81
ratio  :   1.01
```

← dynamic range of x_1

← dynamic range of x_2

← dynamic range of $\hat{x}_1$

← dynamic range of $\hat{x}_2$

Data preprocessing

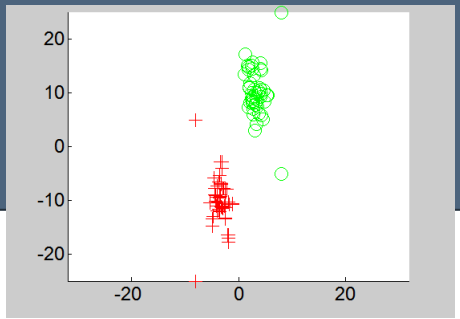

- **Comparison**

mu-sigma normalization | max-min normalization | *softmax* normalization

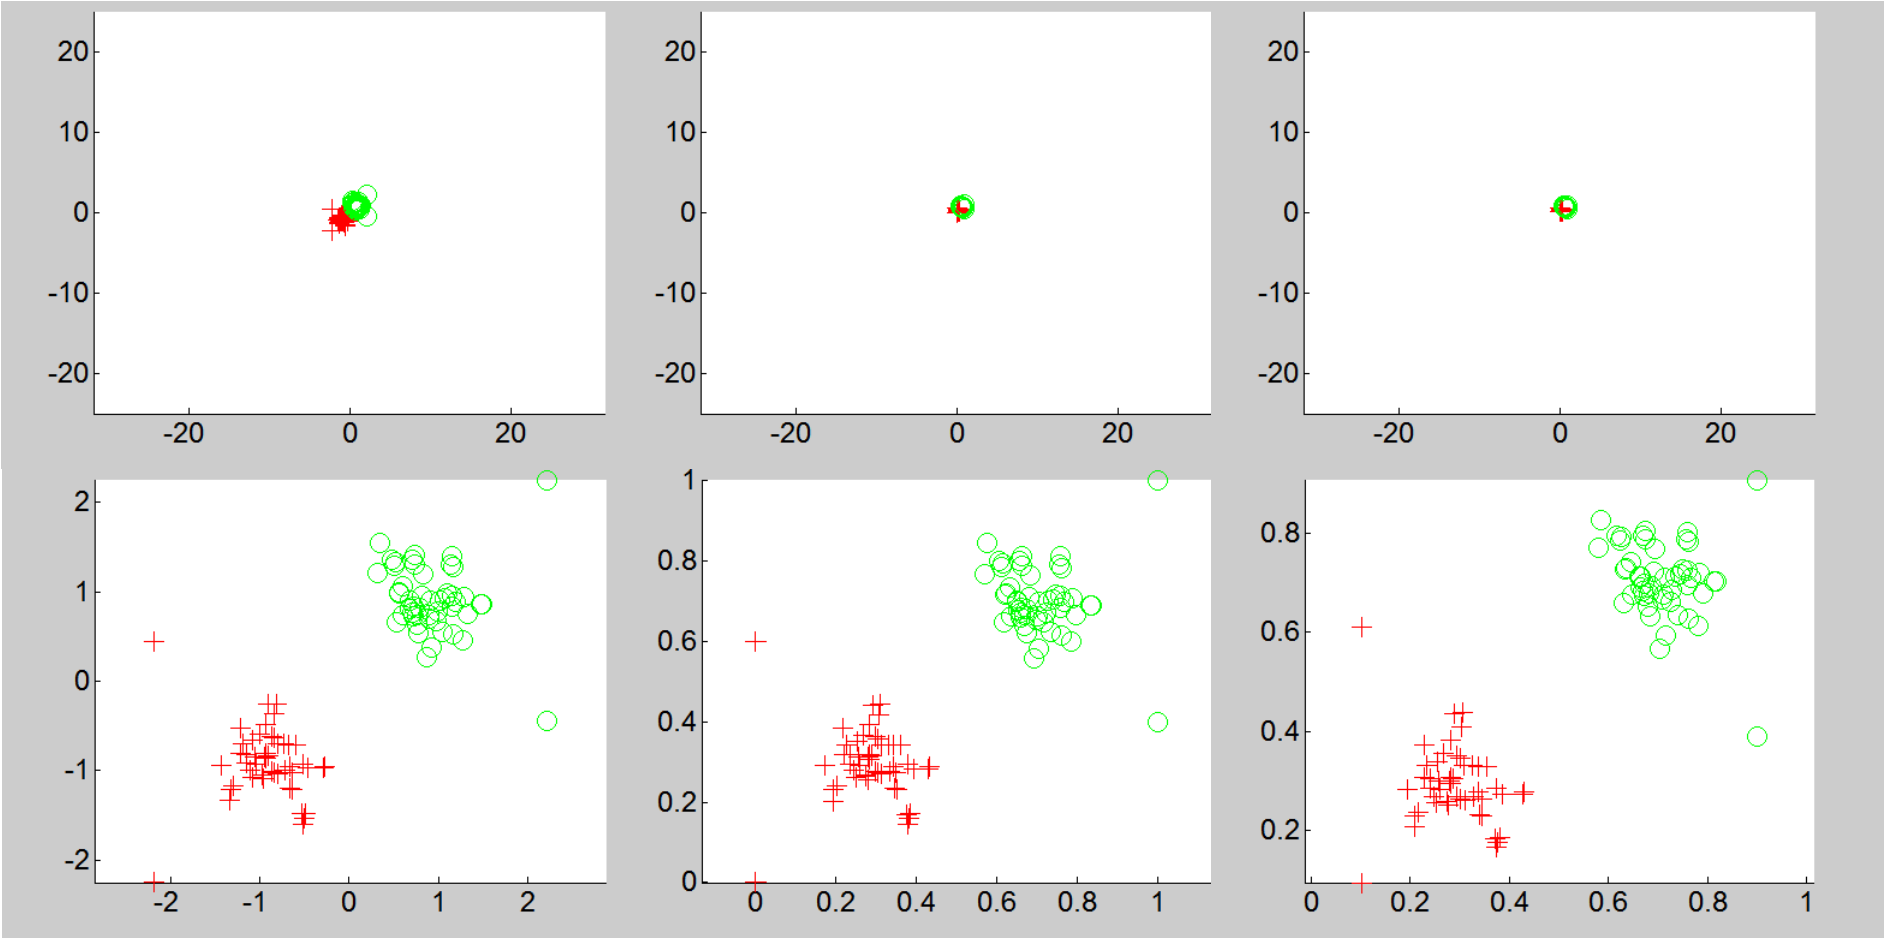

Data preprocessing: Normalization

- Support in **Python**:

```
from sklearn import preprocessing
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
scaler = preprocessing.StandardScaler()
scaler.fit(X)
Xhat = scaler.transform(X)
print(scaler.mean_)
print(scaler.scale_)
print(Xhat.mean(axis=0))
print(Xhat.std(axis=0))
```

```
[5.8433 3.0573 3.7580 1.1993]
[0.8253 0.4344 1.7594 0.7597]
[-1.7e-15 -1.8e-15 -1.7e-15 -1.4e-15]
[1. 1. 1. 1.]
```

$\equiv \mu$

$\equiv \sigma$

```
from sklearn import preprocessing
from sklearn.datasets import load_iris
iris = load_iris()
X = iris.data
scaler = preprocessing.MinMaxScaler()
scaler.fit(X)
Xhat = scaler.transform(X)
print(scaler.min_)
print(scaler.scale_)
print(Xhat.max(axis=0))
print(Xhat.min(axis=0))
```

$\equiv$ – min / (max – min)

```
[-1.1944 -0.8333 -0.1695 -0.0417]
[ 0.2778  0.4167  0.1695  0.4167]
[1. 1. 1. 1.]
[0. 0. 0. 0.]
```

$\equiv$ 1 / (max – min)

- **Filling in missing data**
  - Sometimes, a dataset is incomplete:

x_1	x_2	x_3	x_4
1	2	-4	3
3	2	?	2
2	5	3	?
2.5	?	2	3
1	1	0	2
6	2	4	1

  - In Python, **?** typically appear as **Nan**, Null or None values
  - Machine learning models cannot handle these kind of values
  - Filling with **0s** is not an option

- If the training set is large enough, one can **discard** incomplete samples
- With some datasets, discarding samples is not an option → **heuristic prediction**
  - e.g. fill missing data using the average value from complete samples
  - e.g. fill missing data according to the inherent distribution

Data preprocessing: Missing data

- We will illustrate the process with the *Titanic* dataset:

```
import seaborn as sb
titanic = sb.load_dataset('titanic')
df = titanic.iloc[:,0:13]
print(df.info())
```

```
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column       Non-Null Count  Dtype
---  ------       --------------  -----
 0   survived     891 non-null    int64
 1   pclass       891 non-null    int64
 2   sex          891 non-null    object
 3   age          714 non-null    float64
 4   sibsp        891 non-null    int64
 5   parch        891 non-null    int64
 6   fare         891 non-null    float64
 7   embarked     889 non-null    object
 8   class        891 non-null    category
 9   who          891 non-null    object
 10  adult_male   891 non-null    bool
 11  deck         203 non-null    category
```

  - We can see that column *Age* contains missing (null) values
    - Also other columns, sometimes they are not useful from the ML point of view

  - This can also be obtained by means of the **isnull()** and the **isna()** methods:

```
print(df.isnull().sum())
print(df.isna().sum())
```

```
survived        0
pclass          0
sex             0
age           177
sibsp           0
parch           0
fare            0
embarked        2
class           0
who             0
adult_male      0
deck          688
```

Data preprocessing: Missing data

- We can proceed in several ways:
  - **Delete the columns** with missing data:

```
udf = df.dropna(axis=1)
print(udf.info())
```

```
RangeIndex: 891 entries, 0 to 890
Data columns (total 9 columns):
 #   Column       Non-Null Count  Dtype
---  ------       --------------  -----
 0   survived     891 non-null    int64
 1   pclass       891 non-null    int64
 2   sex          891 non-null    object
 3   sibsp        891 non-null    int64
 4   parch        891 non-null    int64
 5   fare         891 non-null    float64
 6   class        891 non-null    category
 7   who          891 non-null    object
 8   adult_male   891 non-null    bool
```

  - **Delete the rows** with missing data:

```
udf = df.dropna(axis=0)
print(udf.info())
```

    - In this way, we remove too many entries because of the *deck* column:

```
-> 183 entries, 1 to 889
```

    - If we remove first the *deck* column and next the rows with missing data:

```
udf = df.drop('deck', axis=1)
udf.dropna(axis=0,inplace=True)
print(udf.info())
```

```
Int64Index: 712 entries, 0 to 890
Data columns (total 11 columns):
 #   Column       Non-Null Count  Dtype
---  ------       --------------  -----
 0   survived     712 non-null    int64
 1   pclass       712 non-null    int64
 2   sex          712 non-null    object
 3   age          712 non-null    float64
 4   sibsp        712 non-null    int64
 5   parch        712 non-null    int64
 6   fare         712 non-null    float64
 7   embarked     712 non-null    object
 8   class        712 non-null    category
 9   who          712 non-null    object
 10  adult_male   712 non-null    bool
```

Data preprocessing: Missing data

- We can proceed in several ways:
  - Fill the missing values by means of **feature imputation**:

    Numerical data …

    - Fill with the mean
    - Fill with the median
    - Fill with extreme values that do not occur in the data

    Categorical data …

    - Fill with the mode of the distribution
    - Fill with a new label

```
udf = df
udf['age'].fillna(udf['age'].mean(),
                  inplace=True)
```

  - It is good practice to register which data values have been *filled in* before the imputation:

```
udf['missing_age'] = df['age'].isnull()
print(udf.info())
```

```
RangeIndex: 891 entries, 0 to 890
Data columns (total 13 columns):
 #   Column        Non-Null Count  Dtype
---  ------        --------------  -----
 0   survived      891 non-null    int64
 1   pclass        891 non-null    int64
 2   sex           891 non-null    object
 3   age           891 non-null    float64
 4   sibsp         891 non-null    int64
 5   parch         891 non-null    int64
 6   fare          891 non-null    float64
 7   embarked      889 non-null    object
 8   class         891 non-null    category
 9   who           891 non-null    object
 10  adult_male    891 non-null    bool
 11  deck          203 non-null    category
 12  missing_age   891 non-null    bool
```

Data preprocessing: Missing data

- We can proceed in several ways:
  - Fill the missing values using the *SimpleImputer* class for **univariate imputation**

```
from sklearn.impute import SimpleImputer

imp = SimpleImputer(missing_values=np.nan,
                    strategy='mean')

X1 = np.array([[1,       2],
               [np.nan, 3],
               [7,       6]])
imp.fit(X1)

print('X1 = \n', X1)
print('X1* = \n', imp.transform(X1))

X2 = np.array([[np.nan, 2],
               [6, np.nan],
               [7, 6]])
print('X2 = \n', X2)
print('X2* = \n', imp.transform(X2))
```

```
X1 =
 [[ 1.   2.]
  [nan   3.]
  [ 7.   6.]]
X1* =
 [[1. 2.]
  [4. 3.]
  [7. 6.]]
X2 =
 [[nan   2.]
  [ 6. nan]
  [ 7.   6.]]
X2* =
 [[4.    2.  ]
  [6.    3.67]
  [7.    6.  ]]
```

$\overline{X}_1 = [4\,,\,3.67]$

  - **Multivariate imputation** is available experimentally in the *IterativeImputer* class
    - Take into account all columns, instead of only the column with missing values
  - You can also use a **multivariate regression model** to interpolate missing values

Contents

Goodness measures

- **General criterion**:
  - Features should result in a large distance between classes (*between-class distance*) and a reduced variance between class elements (*within-class variance*)
  - **Options**:
    - examine features **in isolation**
      - not optimal, but it serves to discard bad selections easily
    - examine features **in combination**
  - We will consider
    - measures based on **scatter matrices** (examination in combination)
  - but we are not going to consider:
    - measures based on **statistical inference** tests (isolated examination)
      - Each test allows you to work with one feature and two classes only, and requires assumptions about the distribution of classes

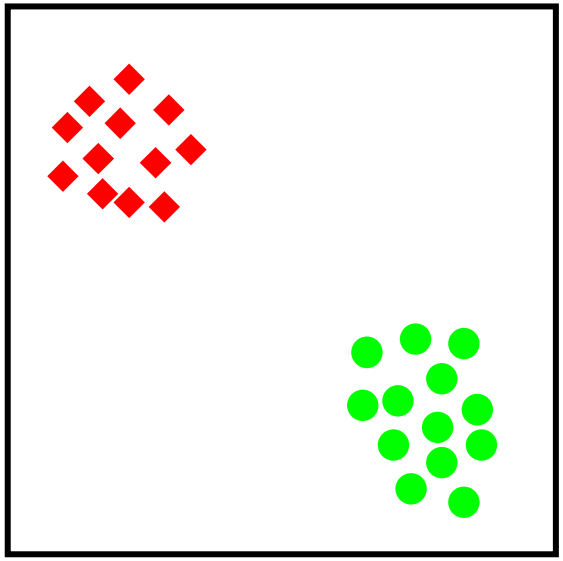

Measures based on scatter matrices

- Measures based on **scatter matrices** allow **multiple classes** and **several characteristics** to be treated simultaneously and do not require the assumption of **normality** in the data
- Let us suppose the following dataset:
  - 10 samples/class, 3 classes, 2 features

	ω_1		ω_2		ω_3	
$x_{k,1}$	7.00	9.03	8.02	7.01	8.99	9.00
$x_{k,2}$	7.05	8.90	7.90	7.13	8.97	9.10
$x_{k,3}$	7.01	8.99	8.12	6.74	9.06	9.29
$x_{k,4}$	7.07	9.19	7.79	7.04	8.88	9.02
$x_{k,5}$	7.03	8.99	7.91	6.88	9.08	8.94
$x_{k,6}$	7.00	8.94	8.26	6.91	9.05	9.17
$x_{k,7}$	6.83	9.05	7.88	7.03	9.08	9.09
$x_{k,8}$	6.96	9.04	7.93	6.89	8.98	9.03
$x_{k,9}$	7.01	8.95	7.86	6.94	8.95	8.98
$x_{k,10}$	7.02	9.07	8.05	6.85	8.99	9.10

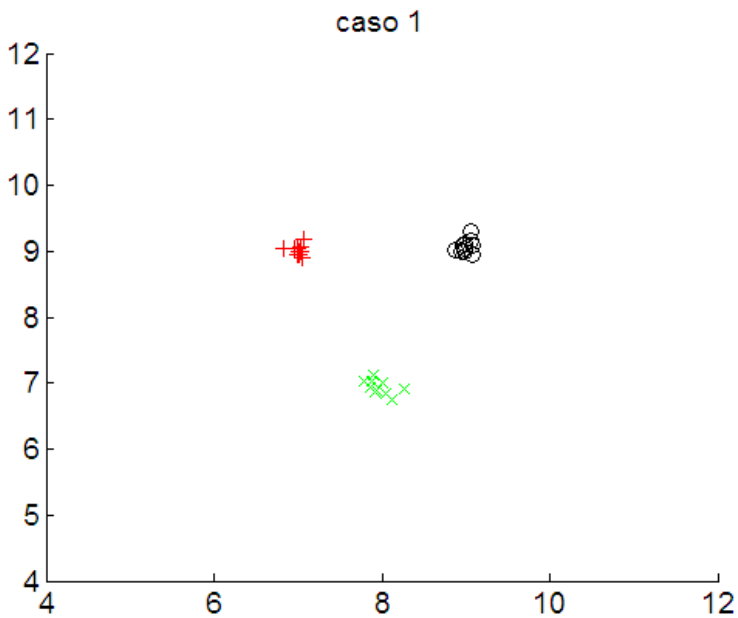

Measures based on scatter matrices

- We first calculate the ***within-class scatter matrix*** (vectors are always column vectors):

$$S_w = \sum_{k=1}^{M} P_k S_k \text{ , where } P_k \approx \frac{n_k}{N} \text{ , } S_k = \frac{1}{n_k - 1} \sum_{j=1}^{n_k} (x_{k,j} - \mu_k)(x_{k,j} - \mu_k)^T$$

where:

$$\mu_k = \frac{1}{n_k} \sum_{j=1}^{n_k} x_{k,j}$$

  - P_k is the **probability *a priori*** of class ω_k
  - S_k is the **covariance matrix** of class ω_k

Measures based on scatter matrices

- Following with the example: M = 3 classes, n_k = 10 samples/class, N = 30 samples

	ω_1		ω_2		ω_3	
$x_{i,1}$	7.00	9.03	8.02	7.01	8.99	9.00
$x_{i,2}$	7.05	8.90	7.90	7.13	8.97	9.10
$x_{i,3}$	7.01	8.99	8.12	6.74	9.06	9.29
$x_{i,4}$	7.07	9.19	7.79	7.04	8.88	9.02
$x_{i,5}$	7.03	8.99	7.91	6.88	9.08	8.94
$x_{i,6}$	7.00	8.94	8.26	6.91	9.05	9.17
$x_{i,7}$	6.83	9.05	7.88	7.03	9.08	9.09
$x_{i,8}$	6.96	9.04	7.93	6.89	8.98	9.03
$x_{i,9}$	7.01	8.95	7.86	6.94	8.95	8.98
$x_{i,10}$	7.02	9.07	8.05	6.85	8.99	9.10
	(7.00, 9.01)		(7.97, 6.94)		(9.00, 9.07)	
	μ_1^T		μ_2^T		μ_3^T	

$$S_1 = \begin{pmatrix} 0.0046 & -0.0001 \\ -0.0001 & 0.0067 \end{pmatrix}$$

$$S_2 = \begin{pmatrix} 0.0201 & -0.0085 \\ -0.0085 & 0.0127 \end{pmatrix}$$

$$S_3 = \begin{pmatrix} 0.0044 & 0.0025 \\ 0.0025 & 0.0104 \end{pmatrix}$$

$$S_w = \begin{pmatrix} 0.0097 & -0.0020 \\ -0.0020 & 0.0100 \end{pmatrix}$$

Measures based on scatter matrices

- We next calculate the ***between-class scatter matrix*** (vectors are always column vectors):

$$S_b = \sum_{k=1}^{M} P_k(\mu_k - \mu_0)(\mu_k - \mu_0)^T \text{, where } P_k \approx \frac{n_k}{N}\text{, } \mu_0 = \sum_{k=1}^{M} P_k \mu_k$$

- Following with the example:

$$\mu_0 = \frac{1}{3}\mu_1 + \frac{1}{3}\mu_2 + \frac{1}{3}\mu_3 = (7.99, 8.34)^T \qquad S_b = \begin{pmatrix} 0.6702 & 0.0321 \\ 0.0321 & 0.9817 \end{pmatrix}$$

- Finally, we obtain the ***mixture scatter matrix***:

$$S_m = S_w + S_b \qquad S_m = \begin{pmatrix} 0.6799 & 0.0301 \\ 0.0301 & 0.9916 \end{pmatrix}$$

- Important properties:
  - **trace(S_w)** measures the dispersion of features inside the classes
  - **trace(S_b)** measures the dispersion of the class centers amongst them
    - ≡ All this permits comparing different features sets among them, i.e. **they are not absolute measures**, but relative

Measures based on scatter matrices

- We can define the following measures:

$$J_1 = \frac{\mathrm{trace}(S_m)}{\mathrm{trace}(S_w)}$$

$$J_2 = \frac{|S_m|}{|S_w|} = |S_w^{-1} S_m|$$

$$J_3 = \mathrm{trace}(S_w^{-1} S_m)$$

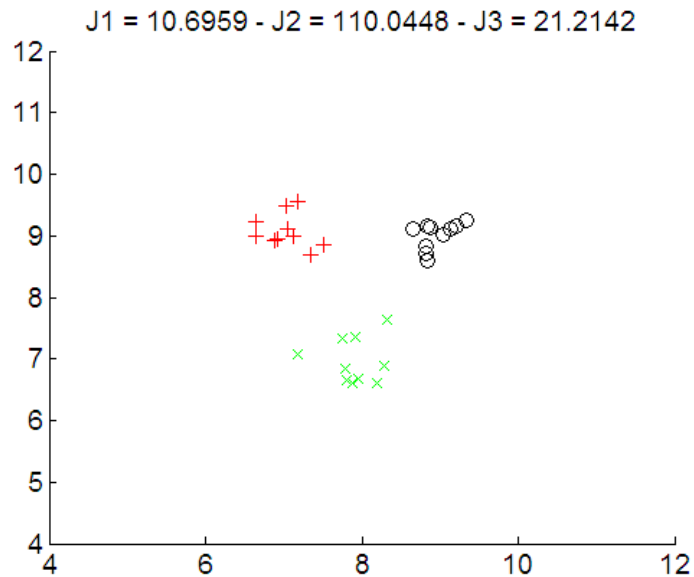

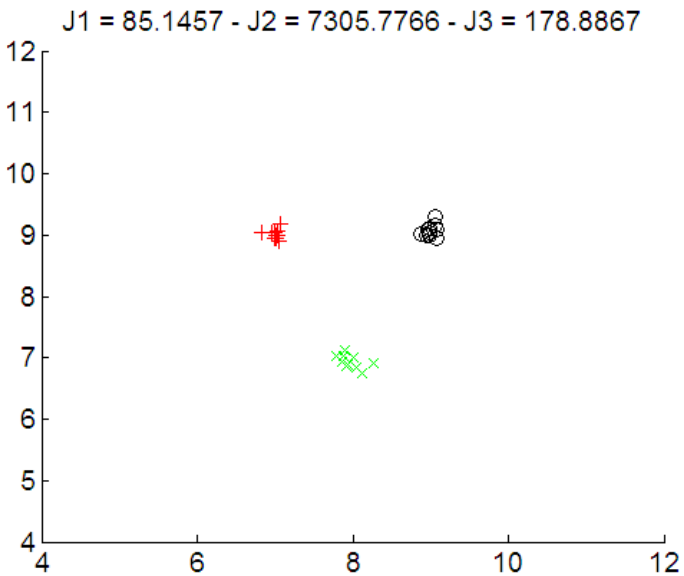

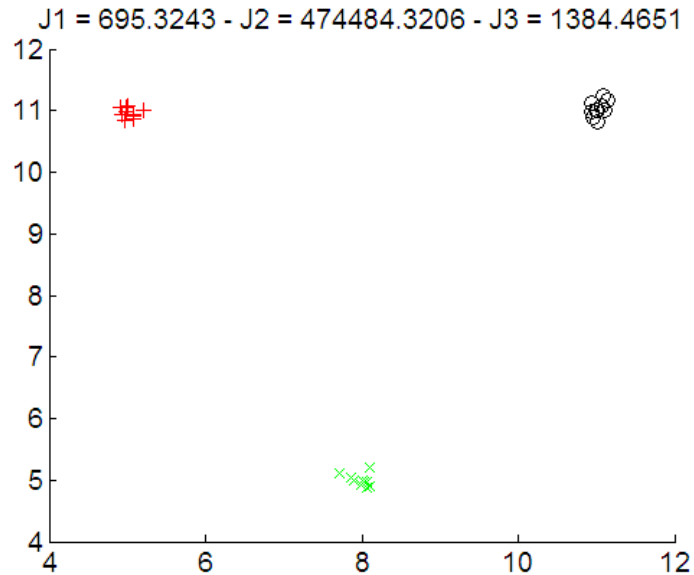

Measures based on scatter matrices

- **1D case** (1 feature) and **2 equiprobable classes**: (= feature by feature and every 2 classes)
  - S_w gets reduced to ${\sigma_1}^2 + {\sigma_2}^2$
  - S_b can be shown to be $\frac{1}{2}(\mu_1 - \mu_2)^2$
- Following with this reasoning, we obtain the *Fisher's Discriminant Ratio* (FDR) for feature f and classes ω_1 and ω_2:

$$\mathrm{FDR}_f(\omega_1, \omega_2) = \frac{(\mu_1 - \mu_2)^2}{\sigma_1^2 + \sigma_2^2}$$

  - FDR can be used to quantify the separability capacity of individual features
  - Similar to the **q-statistics** (measures based on statistical hypothesis testing), but the FDR does not depend on the statistical distribution of the data !!

- Multiclass case, one feature f:

$$\mathrm{FDR}_f^* = \sum_{k_1 \neq k_2} \mathrm{FDR}_f(\omega_{k_1}, \omega_{k_2}), \quad \mathrm{FDR}_f^+ = \min_{k_1 \neq k_2} \{\mathrm{FDR}_f(\omega_{k_1}, \omega_{k_2})\}$$

Measures based on scatter matrices

- For the previous example:

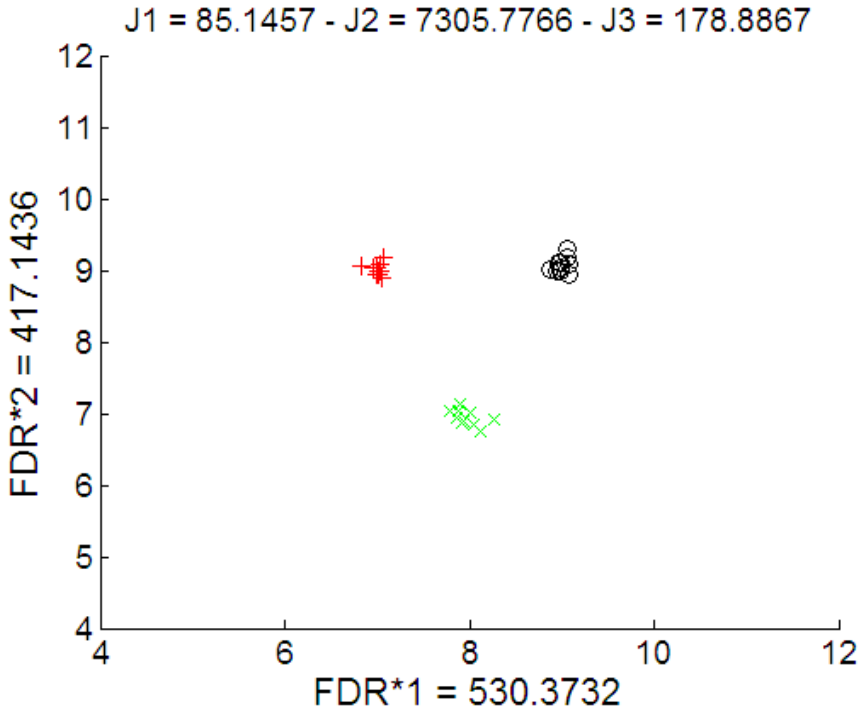

Σ_f FDR$_f$ *	Σ_f FDR$_f$ +
947,5168	417,1436
85,8043	41,7067
7294,3245	3498,2874

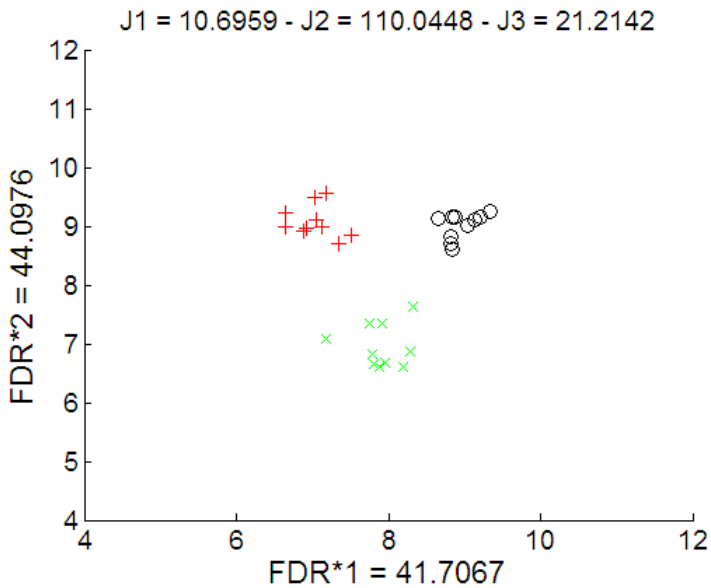

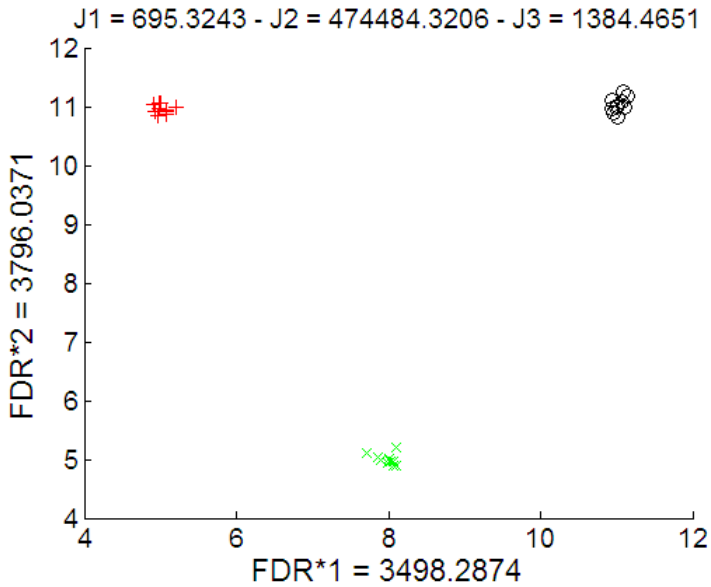

Contents

Feature selection

- Given **C** features in total, this point is about selecting **L** as the most suitable subset
- Several approaches:
  - **Isolated** feature selection
  - **Joint** features selection

- **Isolated selection**

  Sort the **C** features according to a measure of goodness and select the best **L**

```
import matplotlib.pyplot as plt
from sklearn.feature_selection import SelectKBest, f_classif

iris = load_iris()
X, y = iris.data, iris.target
cols = [x[:12] for x in iris.feature_names]

fsel1 = SelectKBest(f_classif, k=4)
X_1 = fsel1.fit_transform(X, y)

sscores = -np.sort(-fsel1.scores_)
ndx = np.argsort(-fsel1.scores_)
scols = [cols[k] for k in ndx]

plt.figure()
plt.bar(scols, sscores)
plt.show()

fsel2 = SelectKBest(f_classif, k=2)
X_2 = fsel2.fit_transform(X, y)
print(X_2.shape)
```

- **F-test** captures linear relationships between features X[:,k] and labels y
  - A highly correlated feature is given a higher score
- Other tests available

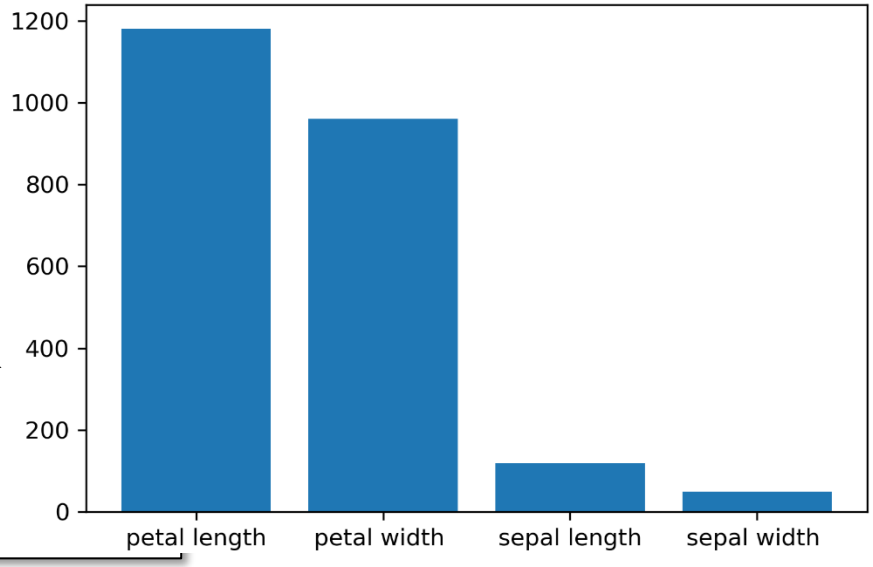


```
(150, 2)
```

- ***GenericUnivariateSelect***`(score_func=score, …)`

- **Joint selection**

  Consider **subsets of L** features

  1) Joint exhaustive selection: go through all combinations and select the best one

  – For example, let us suppose **C = 4** and **L = 2**:

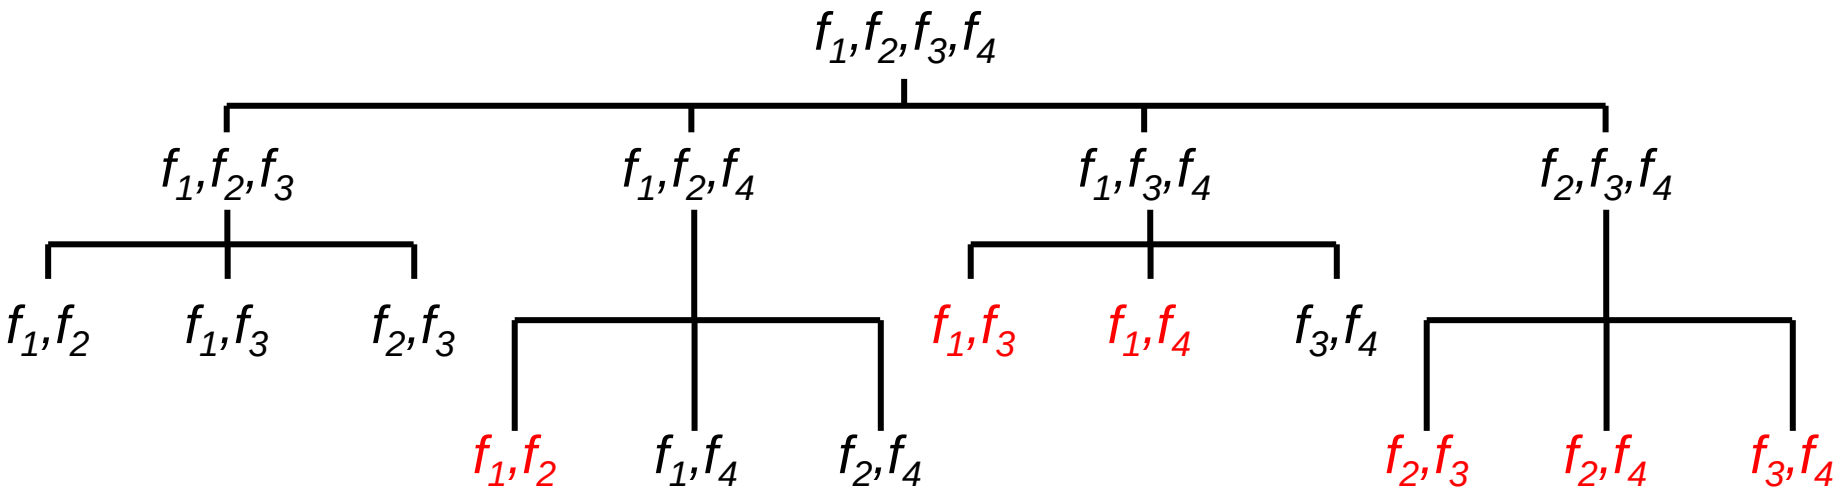

  – 6 combinations in total in this case

  – General case:

$$\binom{C}{L} = \frac{C!}{L!(C-L)!} \quad \text{combinations} \quad \left(\text{e.g. } \binom{20}{5} = 15504\right)$$

2) **Joint suboptimal selection**
   - Go through a subset of combinations
   - It does not guarantee to find the optimal selection but can provide an acceptable selection in less time
   - Two variations:
     - **Backward sequential selection**
     - **Forward sequential selection**

- Suboptimal solutions: **Backward sequential selection**
  - Select a suitable score S, e.g. separability indices J_1, J_2 or J_3, etc.
    - Let us assume S gets higher as better is the combination
  - Starting from **all the features**, progressively remove features until reaching the required amount (of features)
    - At each step keep the combination with the largest score

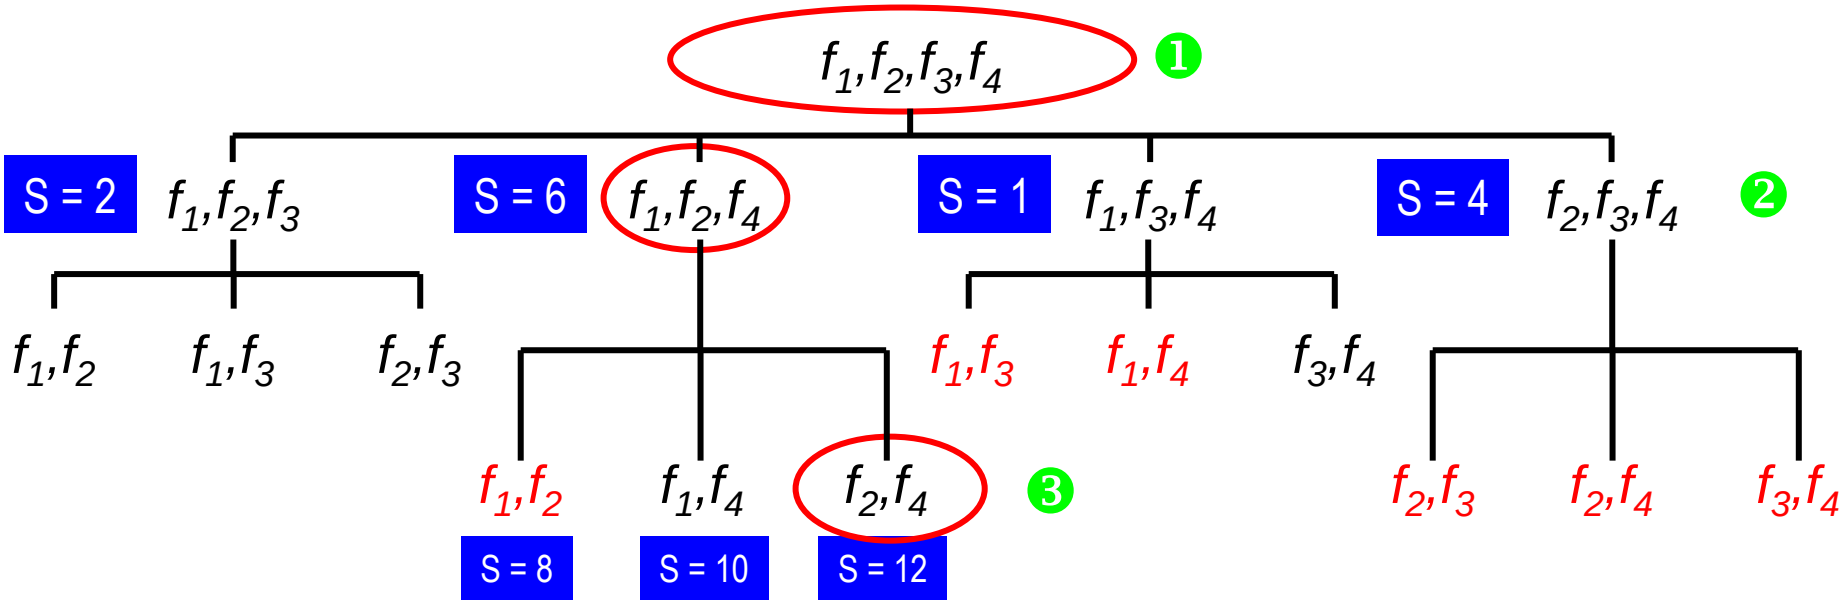

- Suboptimal solutions: **Forward sequential selection**
  - Select a suitable score S, e.g. separability indices J_1, J_2 or J_3, etc.
    - Let us assume S gets higher as better is the combination
  - Starting with **one feature**, progressively add characteristics until the required number of features is reached
    - At each step keep the combination with the largest score

S=1	f_1	S=10	f_2,f_1	S=11	f_2, f_1,f_3
S=5	f_2	S=7	f_2,f_3	S=15	f_2,f_1,f_4
S=2	f_3	S=8	f_2,f_4		
S=3	f_4				
	1		2		3

- Example:

```
import numpy as np
from sklearn.datasets import load_diabetes

diabetes = load_diabetes()
X, y = diabetes.data, diabetes.target
print(diabetes.DESCR)

from sklearn.feature_selection import SequentialFeatureSelector
from sklearn.linear_model import RidgeCV

ridge = RidgeCV(alphas=np.logspace(-6, 6, num=5)).fit(X, y)

sfs_forward = SequentialFeatureSelector(
    ridge, n_features_to_select=2, direction="forward"
).fit(X, y)

sfs_backward = SequentialFeatureSelector(
    ridge, n_features_to_select=2, direction="backward"
).fit(X, y)

feature_names = np.array(diabetes.feature_names)
print(
    "Features selected by forward sequential selection: "
    f"{feature_names[sfs_forward.get_support()]}"
)
print(
    "Features selected by backward sequential selection: "
    f"{feature_names[sfs_backward.get_support()]}"
)
```

```
Diabetes dataset
----------------

Ten baseline variables, age, sex, body mass
index, average blood pressure, and six blood
serum measurements were obtained for each of
n = 442 diabetes patients, as well as the
response of interest, a quantitative measure
of disease progression one year after
baseline.

:Number of Attributes: First 10 columns are
numeric predictive values

:Target: Column 11 is a quantitative measure
of disease progression one year after
baseline

:Attribute Information:
- age age in years
- sex
- bmi body mass index
- bp  average blood pressure
- s1  tc, total serum cholesterol
- s2  ldl, low-density lipoproteins
- s3  hdl, high-density lipoproteins
- s4  tch, total cholesterol / HDL
- s5  ltg, log of serum triglycerides level
- s6  glu, blood sugar level
```

```
Features selected by forward sequential
selection: ['bmi' 's5']

Features selected by backward sequential
selection: ['bmi' 's5']
```

Contents

- Introduction
- Data exploration
- Data preprocessing
- Goodness measures
- Feature selection
- Dimensionality reduction
- Pipelines

Next week, practical session on Thursday: L2

- Laptops available ?
- Description of work to do available on Monday in the webpage
- Revise the notes of lecture 2, try to reproduce the examples

Change in first exam date: L1, L2 - from 19/10 to 2/11/2023
(L3, L4, L6a – 21/12/2023 & L5, L6b – 1/02/2024)

Dimensionality reduction

- **Dimensionality reduction** (DR) refers to the transformation of the original data into a reduced-dimension space, i.e. a new set of variables leading to a reduced set of features
  - Reduce **redundancy** in the feature set, e.g. avoid correlated features
  - Provide a **relevant set of features** for a classifier, aiming at improved performance
  - Give rise to **new meaningful variables** underlying the data, leading to greater understanding of the dataset
- Because of some of the byproducts of DR, it is also termed as **feature extraction**
- One can find several DR methods in the literature:
  - Principal Component Analysis (PCA) and variants (Sparse PCA, Kernel PCA, etc.)
  - Other matrix factorizations:
    - Non-negative Matrix Factorization (NMF)
    - Independent Component Analysis (ICA)
    - Truncated Singular Value Decomposition
  - Multi-dimensional Scaling (MDS)
  - t-distributed Stochastic Neighbor Embedding (t-SNE) – rather for visualizing high-dimensional data

Dimensionality reduction: PCA

- PCA is a popular technique for **analyzing large high-dimensional datasets** to favour the interpretability of the data while **preserving the maximum amount of information**, and maybe enabling the visualization of multidimensional data
  - The aim is to **derive new variables in decreasing order of importance** that are **linear combinations of the original variables** and that are **uncorrelated**
    - The data is linearly **transformed into a new coordinate system** where **most of the variation in the data** can be described with fewer dimensions than the initial data
    - The principal components are a **sequence of orthogonal unit vectors**, i.e. an **orthonormal basis** in which the different individual dimensions of the data are linearly uncorrelated
  - It is a **data-directed** technique: it makes no assumptions of the existence of any kind of property in the data, e.g. clusters
  - Invented in 1901 by Karl Pearson and developed by Harold Hotellin in the 1930s
  - Also named as the **discrete Karhunen-Loeve transform** (KLT) in signal processing, **proper orthogonal decomposition** (POD) in mechanical engineering, etc.

- A simple example:

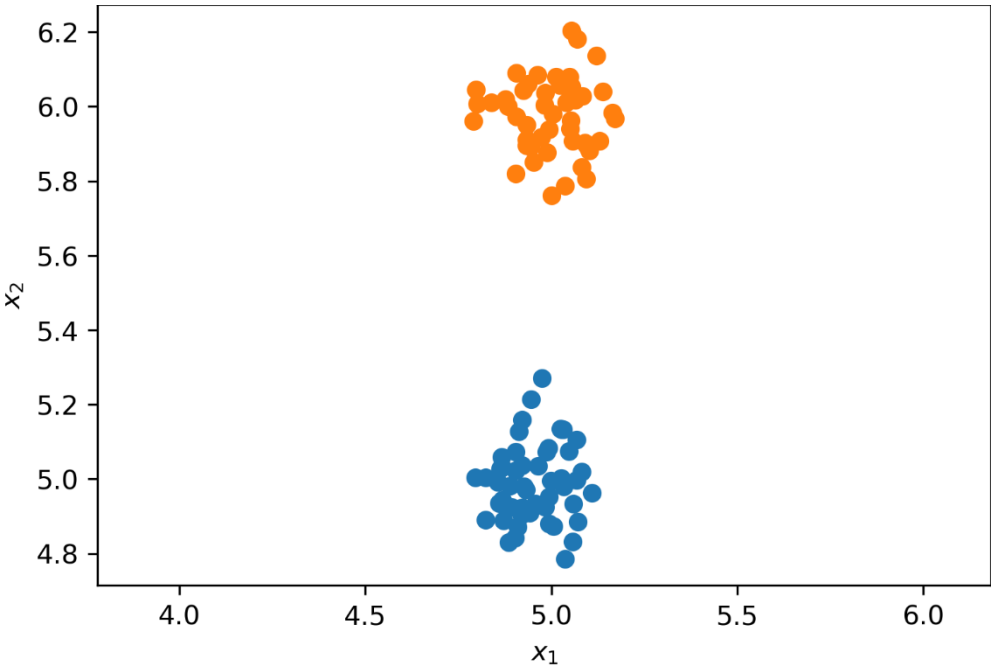

$$\sigma_{x_1}^2 = \frac{1}{N}\sum_{i=1}^{N}(x_{i1}-\mu_1)^2 = 0.008$$

$$\sigma_{x_2}^2 = \frac{1}{N}\sum_{i=1}^{N}(x_{i2}-\mu_2)^2 = 0.256$$

$$\mathrm{cov}(X) = \frac{1}{N}\sum_{i=1}^{N}(x_i-\mu)(x_i-\mu)^T = \begin{pmatrix} 0.008 & 0.012 \\ 0.012 & 0.256 \end{pmatrix}$$

- Which feature should we get rid of?
  - x_1 is not useful from the discrimination point of view
  - x_2 allows discriminating between the two classes

  $\Rightarrow$ x_2 carries more information than x_1 and besides $\sigma_{x1}^2 < \sigma_{x2}^2$

  $\Rightarrow$ if we have to choose, better to get rid of x_1, the one with lowest variance

Dimensionality reduction: PCA

- A more complex example:

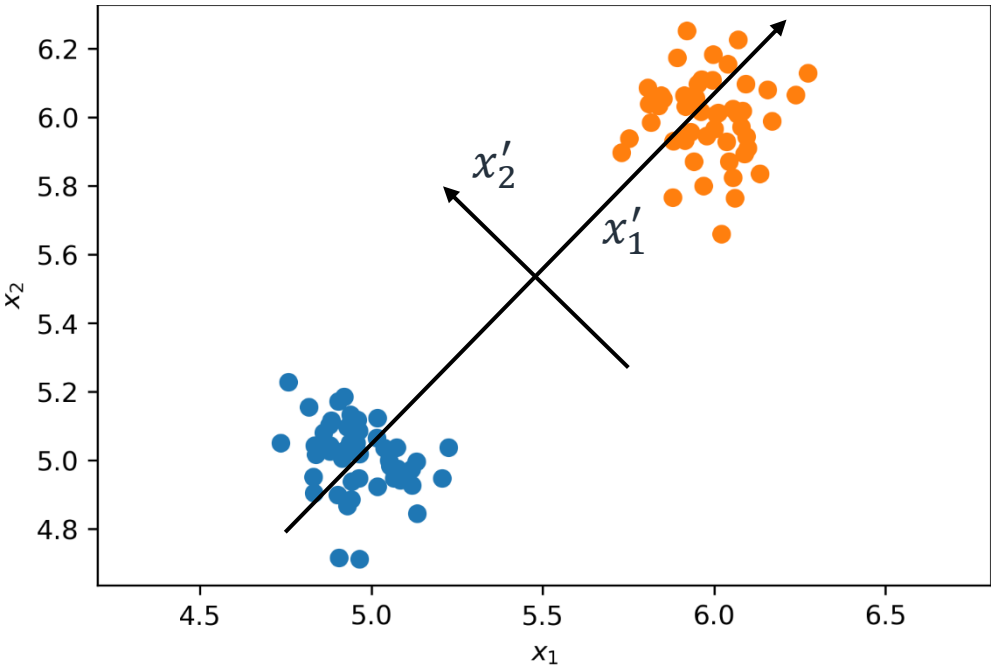

$$\sigma_{x_1}^2 = \frac{1}{N}\sum_{i=1}^{N}(x_{i1}-\mu_1)^2 = 0.280$$

$$\sigma_{x_2}^2 = \frac{1}{N}\sum_{i=1}^{N}(x_{i2}-\mu_2)^2 = 0.257$$

$$\mathrm{cov}(X) = \frac{1}{N}\sum_{i=1}^{N}(x_i-\mu)(x_i-\mu)^T = \begin{pmatrix} 0.280 & 0.254 \\ 0.254 & 0.257 \end{pmatrix}$$

- Which feature should we get rid of?
  - now is not so simple
  - better to consider a different set of features, x_1' and x_2', to have more or less the same separation between classes (besides, we would see more or less the same values as for the previous example in the covariance matrix)

Dimensionality reduction: PCA

- PCA can be thought of as fitting an **L-dimensional hyperellipsoid** to the data
  - Each axis of the ellipsoid represents a principal component
  - If some axis of the ellipsoid is short, it is because the variance along that axis is small
- x_1 and x_2 features are linearly correlated (when x_1 grows, x_2 grows proportionally)
  - but data points projected onto the resulting orthogonal basis are no longer correlated
- To find the axes of the ellipsoid:
  - **center** the values of each feature by subtracting the mean
  - compute the **scatter** or the **covariance matrix**
  - calculate the **eigenvalues** and **eigenvectors** of the resulting matrix
  - the eigenvectors constitute the orthonormal basis (= ellipsoid axes) and each eigenvalue is the variance along the respective axis, i.e. eigenvector

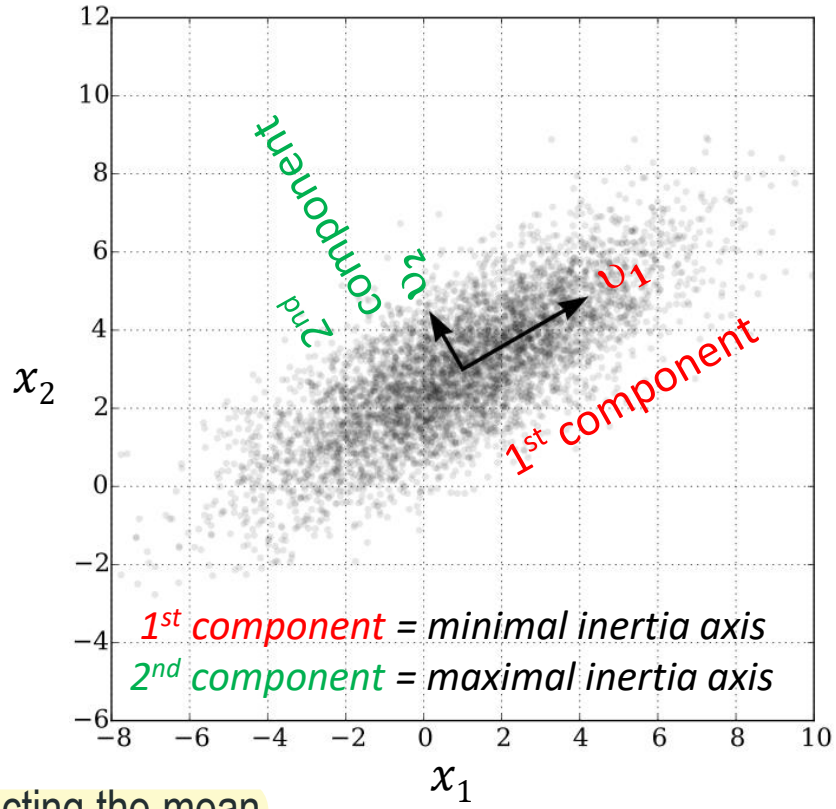

Dimensionality reduction: PCA

- Given the data matrix X which has been **mean-centered** ($X = X_u - \bar{X}$), whose rows contain the data samples $\boldsymbol{x}_i$ and its columns are the feature values, we are looking for a set of vectors υ_i that constitute an **orthonormal basis** where the data is going to be expressed in:

$$t_{i,k} = \mathbf{x}_i \cdot \nu_k \,, \quad i = 1, \ldots, N \text{ and axes } k = 1, \ldots, L$$

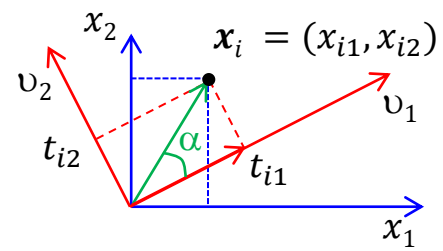

$$\mathbf{x}_i \cdot \nu_1 = \|\nu_1\| \|x_i\| \cos\alpha$$
$$= \|x_i\| \cos\alpha = t_{i,1}$$

- In order to maximize the variance along the first component we look for vector υ_1 such that:

$$\nu_1 = \arg\max_{\|\nu\|=1} \left\{ \sum_i t_{i,\nu}^2 \right\} = \arg\max_{\|\nu\|=1} \left\{ \sum_i (\mathbf{x}_i \cdot \nu)^2 \right\}$$

- In matrix form, this becomes:

$$\nu_1 = \arg\max_{\|\nu\|=1} \left\{ \|X\nu\|^2 \right\} = \arg\max_{\|\nu\|=1} \left\{ \nu^T X^T X \nu \right\}$$

where $A = X^T X$ is the **scatter matrix** of X_u ($\equiv$ covariance matrix if divided by *N-1*).

Dimensionality reduction: PCA

- To find the constrained maximization problem we build the **Lagrangian function** $L(\upsilon)$ as follows:

$$\max_{\|\nu\|=1} \ \nu^T A\nu \Rightarrow \max \ L(\nu) = \nu^T A\nu - \lambda(\nu^T\nu - 1)$$

- The solution is given by:

$$\frac{\partial L}{\partial \nu} = 2A\nu - 2\lambda\nu = 0 \Rightarrow A\nu = \lambda\nu$$

$$\frac{\partial L}{\partial \lambda} = \nu^T\nu - 1 = 0 \Rightarrow \nu^T\nu = 1$$

- Equation $A\upsilon = \lambda\upsilon$ has L solutions (υ_i, λ_i) for $A_{L\times L}$, i.e. $A\upsilon_i = \lambda_i\upsilon_i \ , \forall i$ which corresponds to the **eigendecomposition** of matrix $A_{L\times L}$, which in matrix form is given by:

$$A = V\,D\,V^{-1}, \text{ with } V = \left[\begin{array}{c|c|c|c} \uparrow & \uparrow & & \uparrow \\ \nu_1 & \nu_2 & \dots & \nu_L \\ \downarrow & \downarrow & & \downarrow \end{array}\right] \text{ and } D = \begin{bmatrix} \lambda_1 & 0 & \dots & 0 \\ 0 & \lambda_2 & \dots & 0 \\ & & \ddots & \\ 0 & 0 & \dots & \lambda_L \end{bmatrix}$$

- Linear algebra libraries typically return **unit eigenvectors**, so that all equations are satisfied, and also ordered from largest eigenvalue to lowest eigenvalue:

$$\lambda_1 \geq \lambda_2 \geq \cdots \geq \lambda_L \geq 0$$

Dimensionality reduction: PCA

- Then, we have to find the eigenvector υ that gives rise to

$$\max\ \nu^T A\nu = \max\ \nu^T V D V^{-1}\nu = \max\ \nu^T V D V^T \nu$$

- Let us now suppose that $\upsilon = \upsilon_j$. Then:

$$\nu^T V = \mathbf{e}_j,\ \text{where } \mathbf{e}_j = \big(0,\ \ldots\ ,\ 0,\ \overset{j\downarrow}{1},\ 0,\ \ldots\ ,\ 0\big)$$

and:

$$\max\ \nu^T V D V^T \nu = \max\ \mathbf{e}_j^T D \mathbf{e}_j = \max\ \lambda_j$$

- The first component of PCA is therefore the eigenvector associated to the **largest eigenvalue**, and so $\upsilon = \upsilon_1$ if the eigenvalues are sorted.

Dimensionality reduction: PCA

- To find the second component, we have to maximize for the remaining variance:

$$\max_{\nu} \ \nu^T A\nu - \nu_1^T A\nu_1$$
$$\text{s.t. } \|\nu\| = 1 \text{ and } \nu^T\nu_1 = 0$$

- Then, the Lagrangian function becomes:

$$L(\nu) = \nu^T A\nu - \nu_1^T A\nu_1 - \lambda\left(\nu^T\nu - 1\right) - \mu\left(\nu^T\nu_1\right)$$

and
$$\frac{\partial L}{\partial \nu} = 2A\nu - 2\lambda\nu - \mu\nu_1 = 0 \Rightarrow A\nu = \lambda\nu$$
$$\frac{\partial L}{\partial \lambda} = \nu^T\nu - 1 = 0 \Rightarrow \nu^T\nu = 1$$
$$\frac{\partial L}{\partial \mu} = \nu^T\nu_1 = 0$$

- Referring to the first equation:

$$\begin{aligned}\nu_1^T\left(2A\nu - 2\lambda\nu - \mu\nu_1\right) &= 2\nu_1^T A\nu - 2\lambda\nu_1^T\nu - \mu\nu_1^T\nu_1 \\ &= 2\nu_1^T A\nu - \mu = 0 \\ &\Rightarrow \mu = 2\nu_1^T A\nu = 2\nu^T A\nu_1 = 2\nu^T\lambda_1\nu_1 = 0\end{aligned}$$

- Therefore, the second component is another eigenvector of A and hence has to be the second eigenvector of A, $\upsilon = \upsilon_2$, since λ_2 is the second largest eigenvalue of A.

Dimensionality reduction: PCA

- The remaining components can be proved to be the remaining eigenvectors, ordered by the corresponding eigenvalue from higher to lower.
- Now that we know that the components are the eigenvectors of matrix $X^T X$, we have to deal with the **reduced-dimension representation**.
- To this end, we consider the **fraction of the total variance** that is accounted for by the first p components:

$$\frac{\sum_i^p \operatorname{var}[\nu_i]}{\sum_i^L \operatorname{var}[\nu_i]} = \frac{\sum_i^p \lambda_i}{\sum_i^L \lambda_i}$$

  - We can specify a threshold τ on this ratio to choose the number of components necessary to account for at least a τ fraction of the total variance.
- We can also plot the eigenvalues in decreasing order (**SCREE plot**) and look for the component for which the accounted variance falls sharply:
  - eigenvalues
  - fraction of total variance: $\frac{\lambda_i}{\sum_j \lambda_j}$

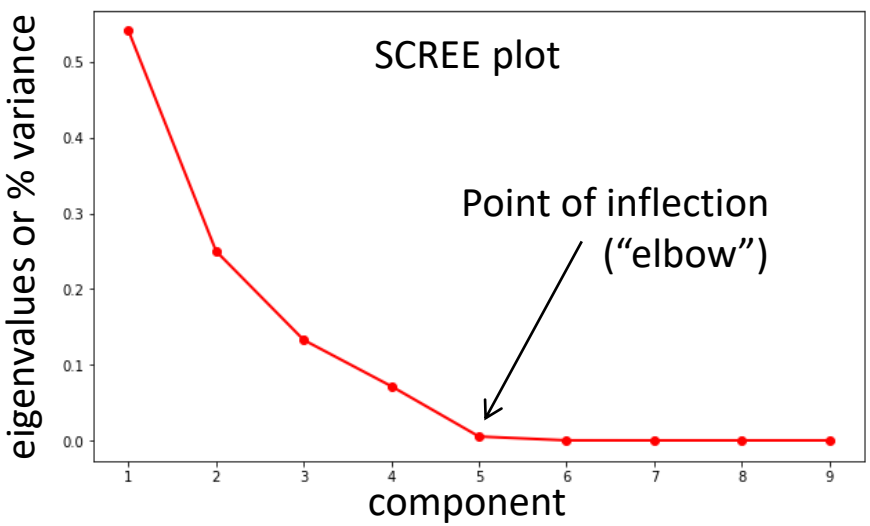

Dimensionality reduction: PCA

- Once we have decided to make use of p components, we can find the reduced-dimensionality vectors/samples: (we do not refer to a particular sample x_i)

centered data	**uncentered data**
$\chi_p = V_{:p}^T \mathbf{x}$	$\chi_p = V_{:p}^T (\mathbf{x} - \mu)$

where χ_p is the reduced-dimension sample, μ is the mean and

$$V_{:p} = \left[\nu_1 \mid \nu_2 \mid \ldots \mid \nu_p\right]_{L \times p}$$

- The dimensionality reduction operation gives rise to an **error of representation**. This makes interesting to know the representation $\boldsymbol{x}_p$ of χ_p in the original L-dimensional space:

	centered data	**uncentered data**
to transformed space	$\chi = V^T \mathbf{x}$	$\chi = V^T (\mathbf{x} - \mu)$
to original space	$\mathbf{x} = V \chi$	$\mathbf{x} = V \chi + \mu$
reduced dimension, but in the original space	$\mathbf{x}_p = V \begin{pmatrix} \chi_p \\ \mathbf{0} \end{pmatrix} = V_{:p} \chi_p = V_{:p} V_{:p}^T \mathbf{x}$	$\mathbf{x}_p = V \begin{pmatrix} \chi_p \\ \mathbf{0} \end{pmatrix} + \mu = V_{:p} \chi_p + \mu = V_{:p} V_{:p}^T (\mathbf{x} - \mu) + \mu$

Dimensionality reduction: PCA

- <u>Example 1</u>:
  Let us consider again the 4-dimensional **Iris dataset**

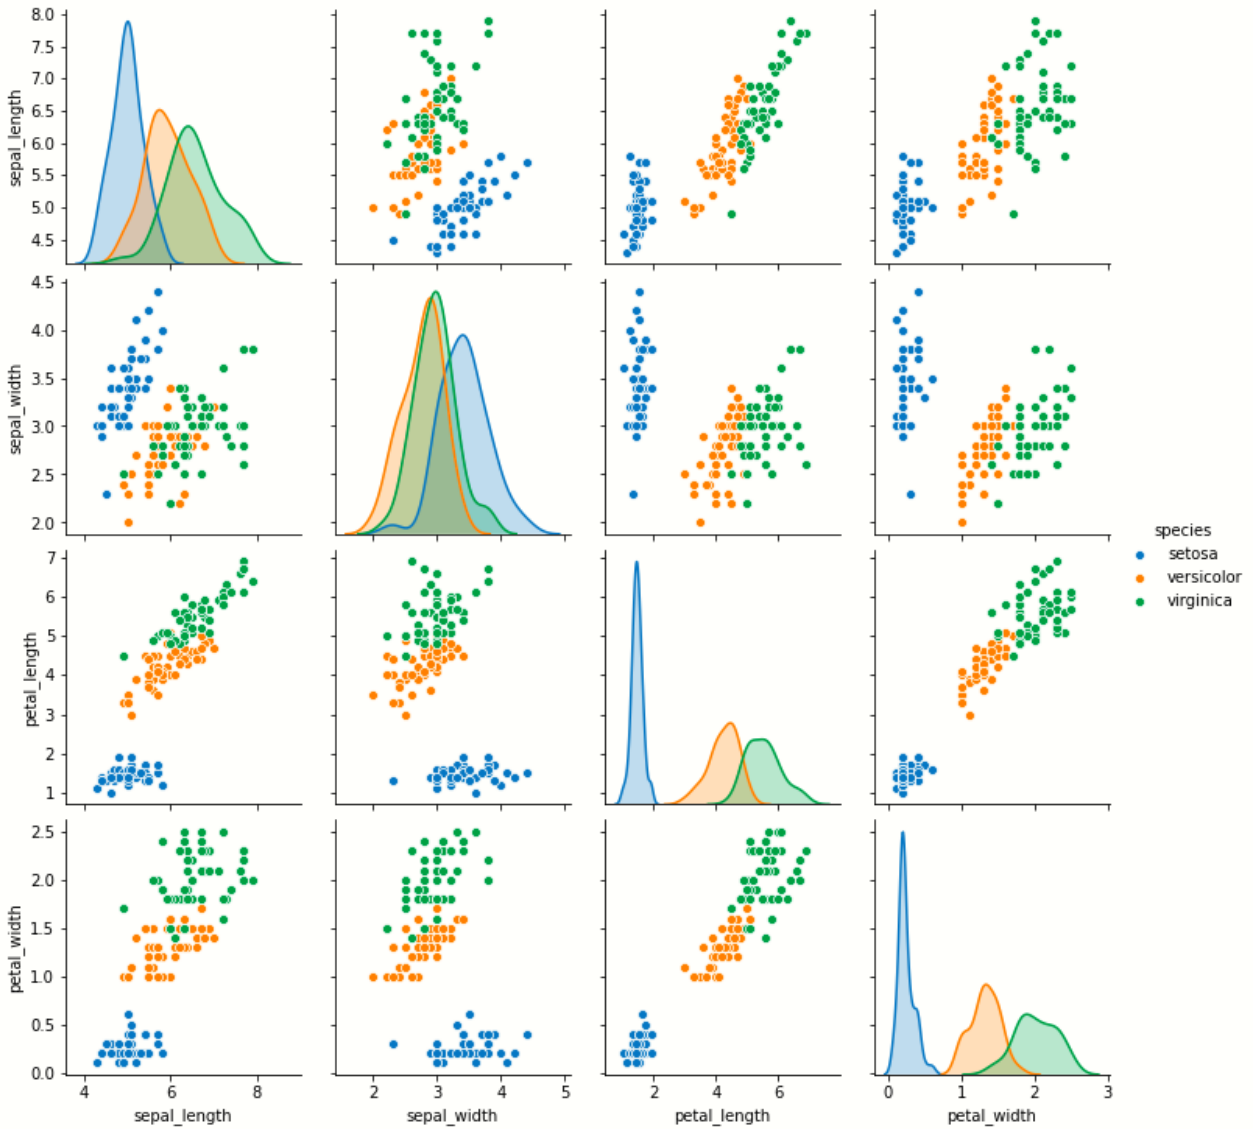

Dimensionality reduction: PCA

- Example 1:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.decomposition import PCA

iris = datasets.load_iris()

X = iris.data
y = iris.target

pca = PCA(n_components=2)
# fit method already includes centering
Xr = pca.fit(X).transform(X)

plt.figure()
for c in range(3):
    i = np.where(y == c)[0]
    plt.scatter(Xr[i,0],Xr[i,1])
plt.show()

pca = PCA(n_components=4)
pca.fit(X)

print(pca.explained_variance_)
print(pca.explained_variance_ratio_)
```

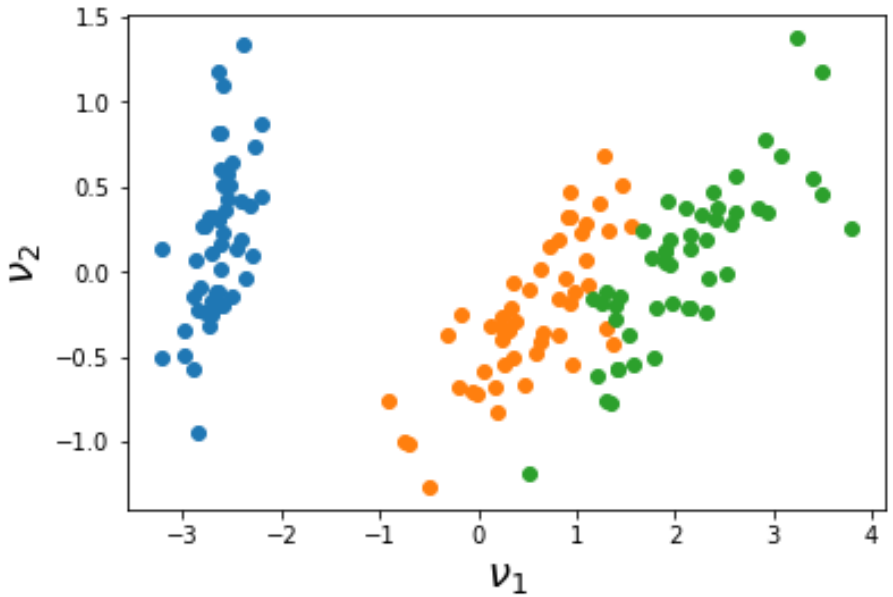


```
[4.2282 0.2427 0.0782 0.0238]
[0.9246 0.0531 0.0171 0.0052]
```

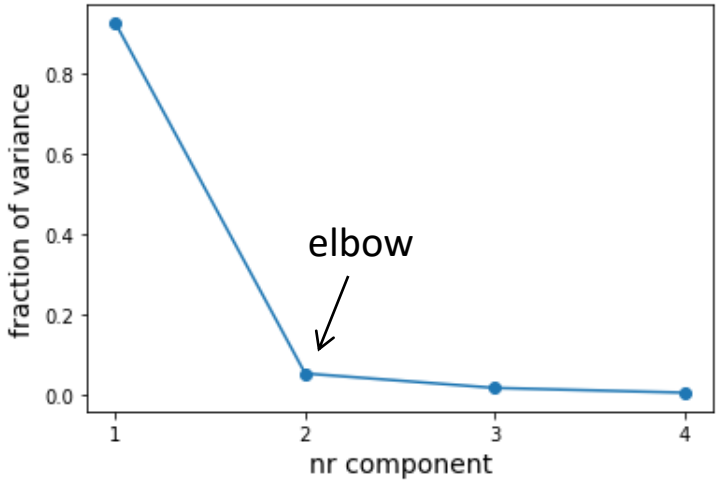

- Example 1:

```
(continued)
pca = PCA(n_components=2)
pca.fit(X)
Xr = pca.transform(X)
X_ = pca.inverse_transform(Xr)

import numpy as np
from math import sqrt
error_matrix = X - X_
error_sq = np.sum(np.sum((error_matrix)**2, axis=1))
error = sqrt(error_sq)
N = X.shape[0]
print('total error = %f, total error/sample = %f' % (error, error / N))
m = np.min(np.abs(error_matrix), axis=0)
print('min. errors: %f %f %f %f' % (m[0], m[1], m[2], m[3]))
m = np.max(np.abs(error_matrix), axis=0)
print('max. errors: %f %f %f %f' % (m[0], m[1], m[2], m[3]))
m = np.min(X, axis=0)
print('min. values: %f %f %f %f' % (m[0], m[1], m[2], m[3]))
m = np.max(X, axis=0)
print('max. values: %f %f %f %f' % (m[0], m[1], m[2], m[3]))
```

```
total error = 3.899313, total error/sample = 0.025995
min. errors: 0.001556 0.001401 0.000492 0.000810
max. errors: 0.451606 0.463801 0.233806 0.591713
min. values: 4.300000 2.000000 1.000000 0.100000
max. values: 7.900000 4.400000 6.900000 2.500000
```

Dimensionality reduction: PCA

- Example 2: **Olivetti faces dataset**,
  400 faces of 64 × 64 pixels,
  4096 dimensions

Faces from dataset

Eigenfaces - PCA using randomized SVD

Transformed faces (10 components)

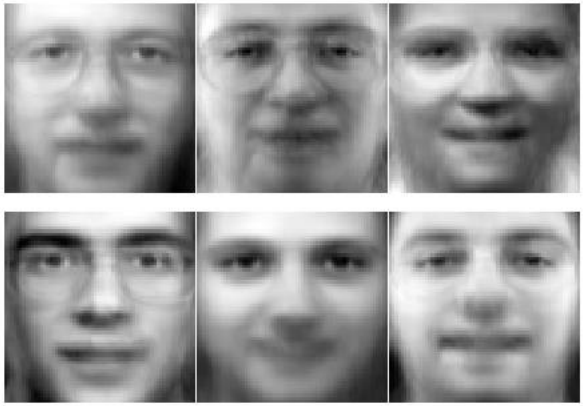

Transformed faces (100 components)

Transformed faces (200 components)

Dimensionality reduction: PCA

- <u>Example 2</u>: **Olivetti faces dataset**,
  400 faces of 64 × 64 pixels,
  4096 dimensions

Faces from dataset

Eigenfaces - PCA using randomized SVD

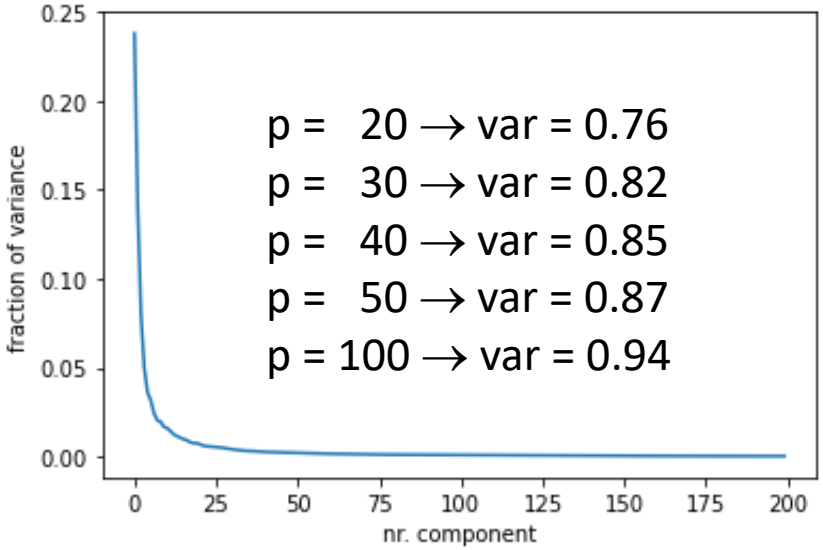

Transformed faces (200 components)

Contents

Pipelines

- Scikit-learn allows chaining steps to transform data until reaching the final estimator:

```
from sklearn import datasets
from sklearn.decomposition import PCA
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

# Define a pipeline to search for the best combination of PCA truncation
scaler = StandardScaler() # mu-sigma scaler to normalize inputs
pca = PCA()               # dimensionality reduction
logistic = LogisticRegression(max_iter=10000, tol=0.1) # classifier
pipe = Pipeline(steps=[("scaler", scaler), ("pca", pca), ("logistic", logistic)])

X_digits, y_digits = datasets.load_digits(return_X_y=True)
# Parameters of pipelines can be set using '__' separated parameter names:
param_grid = {
        "pca__n_components": [10, 20, 30, 40, 50, 60],
        "logistic__C": [0.01, 10, 100],
}
search = GridSearchCV(pipe, param_grid, n_jobs=-1)
search.fit(X_digits, y_digits)
print("Best configuration (CV score=%0.3f):" % search.best_score_)
print(search.best_params_)
```

```
Best configuration (CV score=0.917):
{'logistic__C': 10, 'pca__n_components': 50}
```

Lecture 2:
Data analysis

Universitat de les Illes Balears
Departament de Ciències Matemàtiques i Informàtica

11752 Aprendizaje Automático
11752 Machine Learning
Máster Universitario
en Sistemas Inteligentes

Alberto ORTIZ RODRÍGUEZ